

# *BoilerNet*

## *Final Design Review*

Akshath Raghav Ravikiran, Aneesh Poddutur,  
Gautam Nambiar, Gokul Harikrishnan

04/30/2025



# *Team Introduction*



Gautam Nambiar

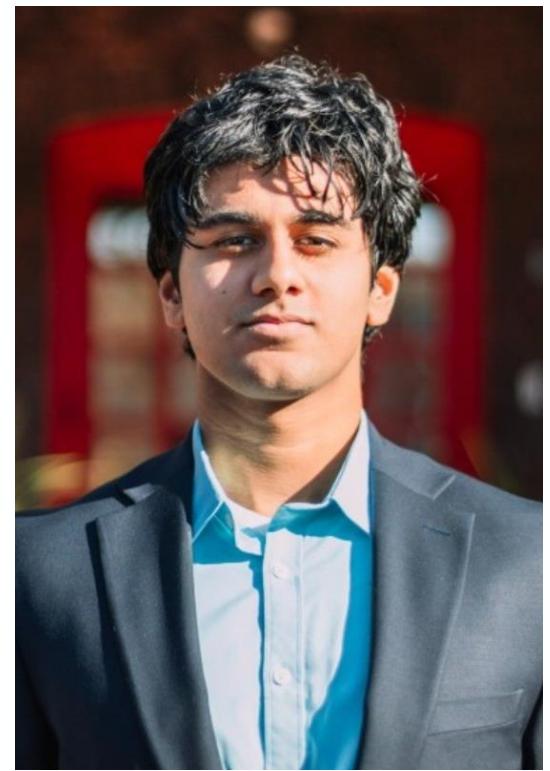
BSCmpE est. May 2025



Akshath Raghav  
Ravikiran  
BSCmpE est. May 2025



Gokul Harikrishnan  
BSCmpE est. May 2025



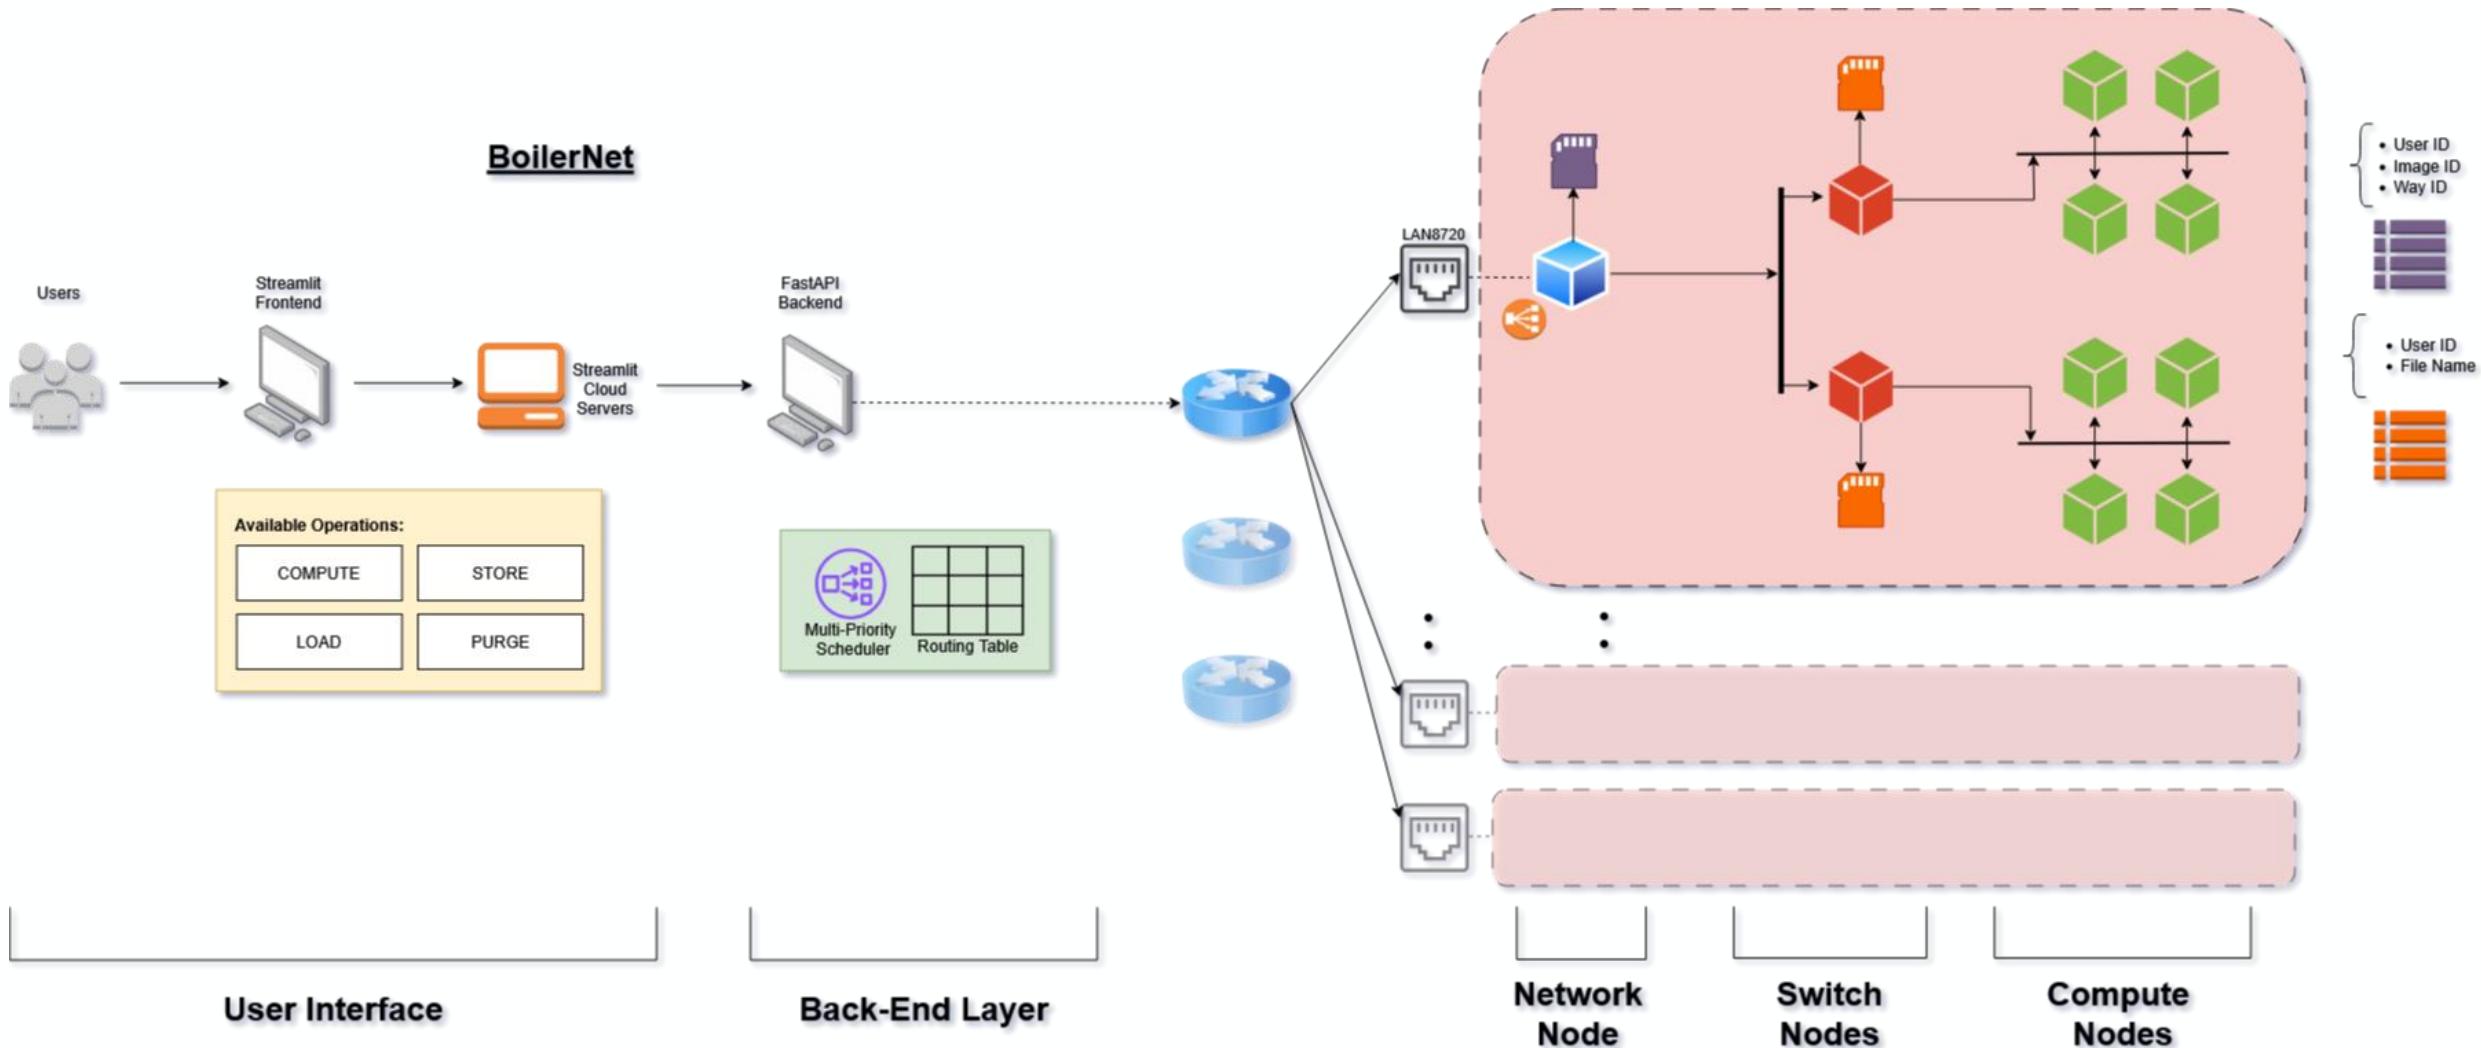
Aneesh Poddutur  
BSEE est. May 2025

# BoilerNet

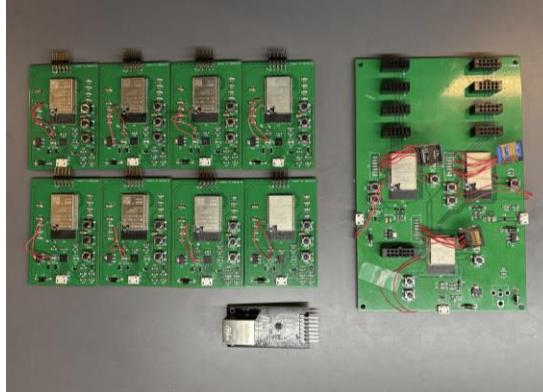
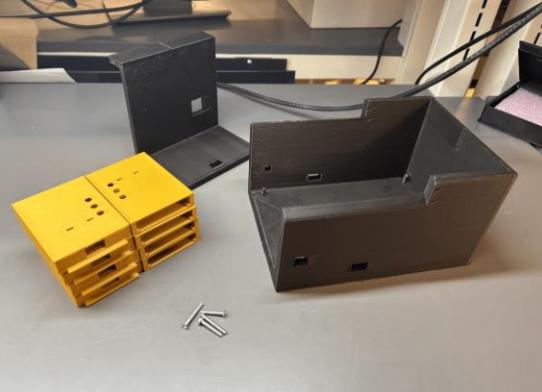
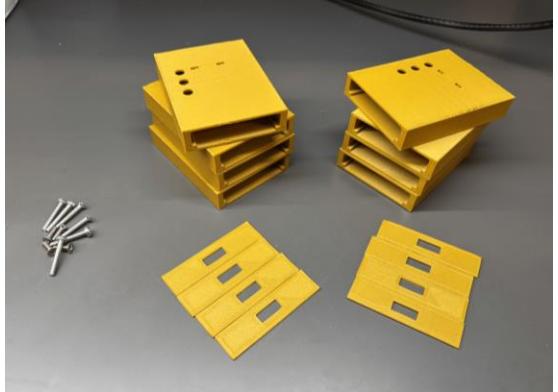
- Portable
- Swappable Compute Workloads
- Low Power
- Scalable
- Throughput-optimized



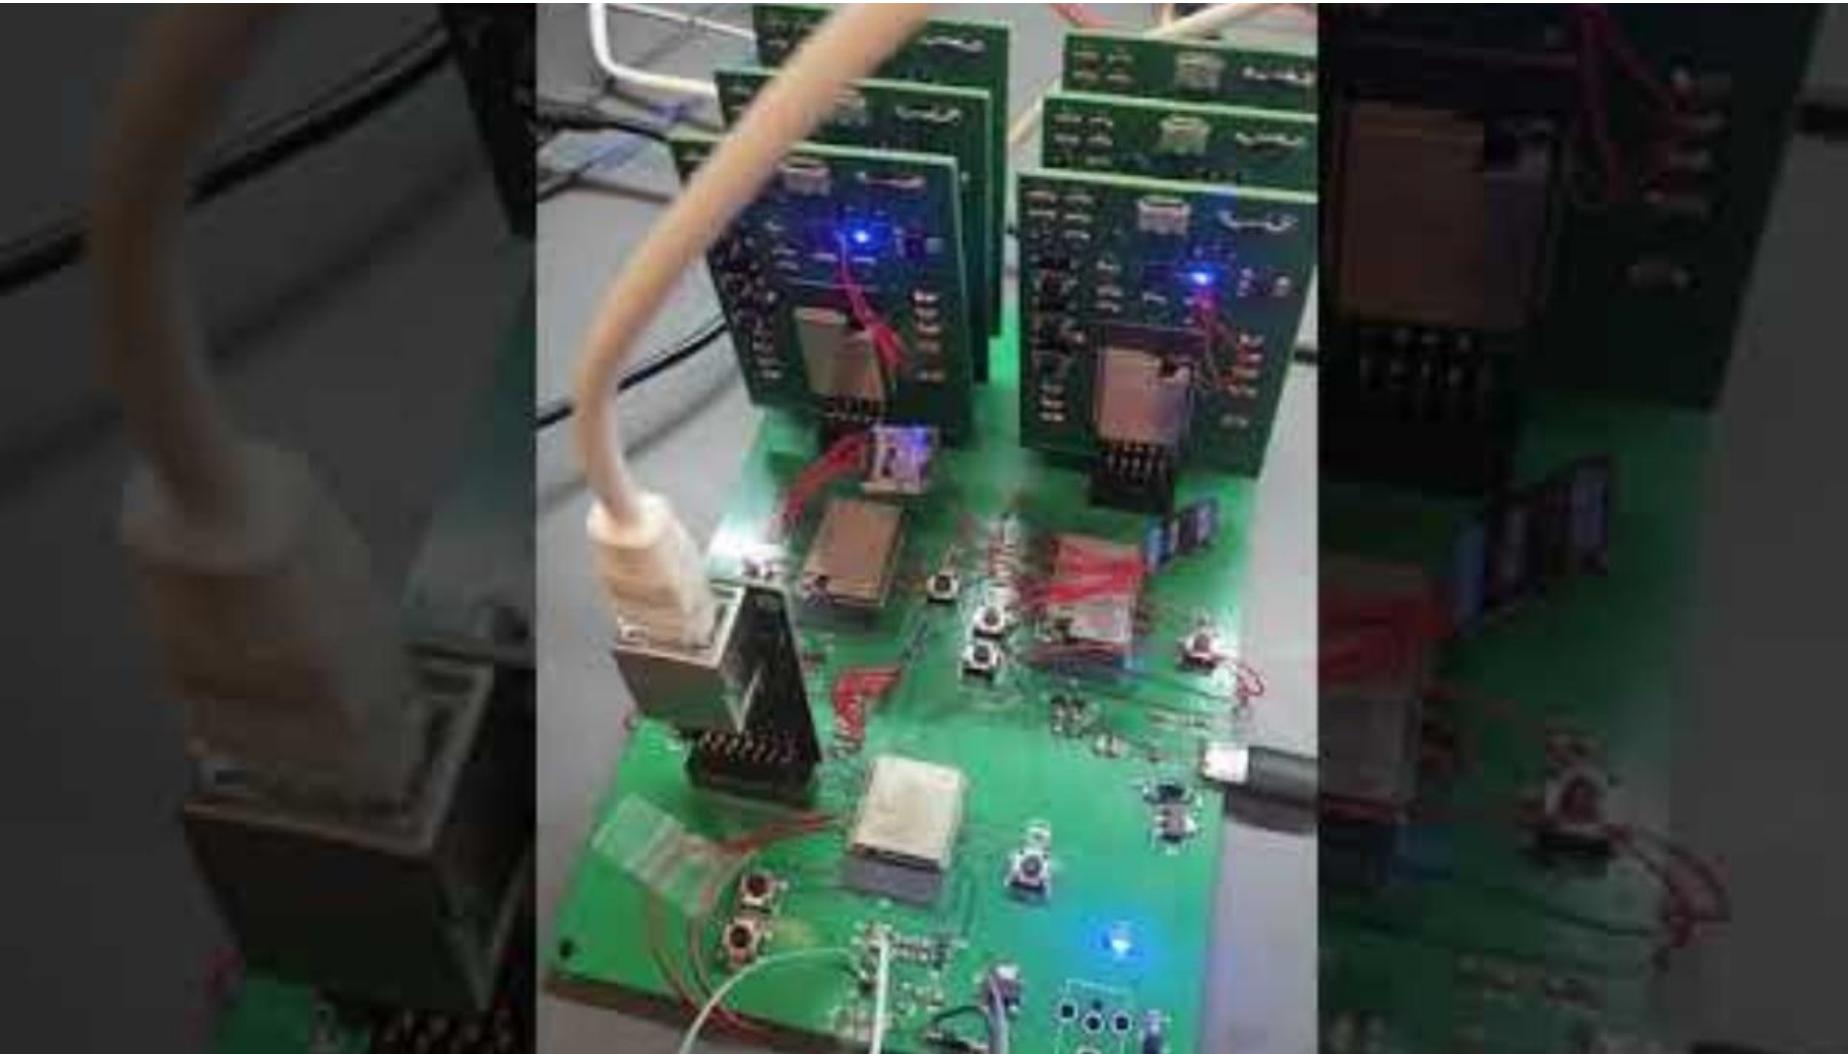
# Block Diagram



# *Mechanical Design*



# *System Integration*



# *System Requirements & Status*

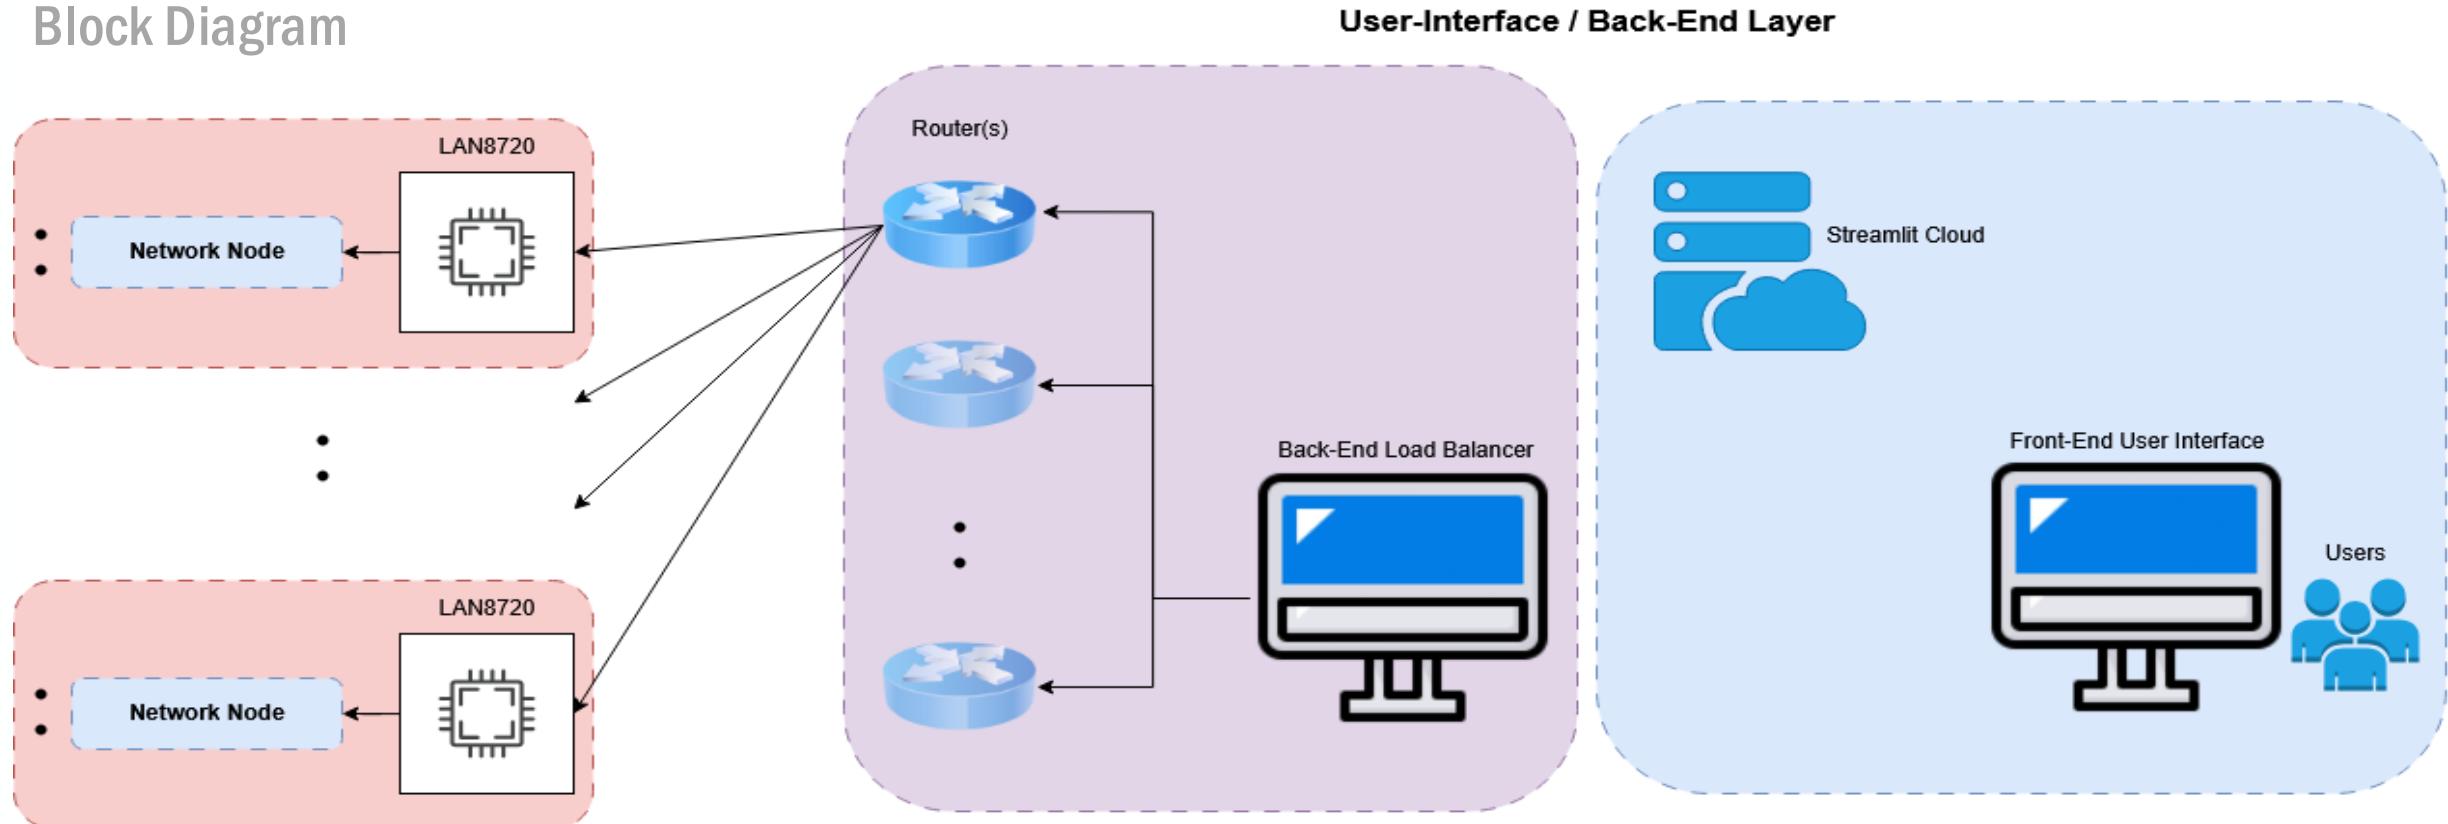
- **Requirement 1:** The User-Interface will allow the User to choose the Operation and define a Job with a (limited) number of Tasks for the cluster to work on. This improves User Experience.
- **Requirement 2:** The Backend will be responsible for buffering Jobs and sending in Tasks serially into the Network Node. It will contain a multi-priority-messaging queue. This ensures no Operation request is starved, and helps in prioritizing I/O heavy operations.
- **Requirement 3:** The Network Node will send the entire task into one of the Ways. Way Selection must happen inside the Network Node. This promotes scalability by decoupling responsibilities.
- **Requirement 4:** The Internal Switch Nodes are responsible for Data Parallelism at the Compute level, and must deal with Compute failures dynamically.

# *System Requirements & Status (Contd.)*

- **Requirement 6:** The Compute Nodes will buffer the PNG-encoded packets as they arrive, and will perform decoding in-memory at runtime. This ensures parallelism can be abstracted away to switches.
- **Requirement 7:** The Compute Nodes are to be developed onto Cartidge-like PCBs, allowing for slotting and unslotting with ease. This allows for swappable workloads.
- **Requirement 8:** The Compute Cluster, as a whole, must be contained within it's 3-D printed enclosure, made in Black and Gold colors. This ensures portability.

# User-Interface & Back-End Layer

## Block Diagram



### User-Interface Highlights:

- Supports multi-file upload per Job with task-level operation selection
- Polls Backend for task status and auto-refreshes UI on completion
- Enables repeated STORE, COMPUTE, LOAD, PURGE on any Task

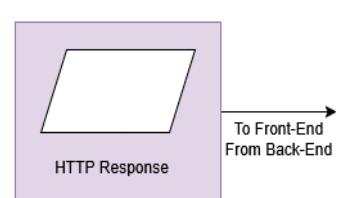
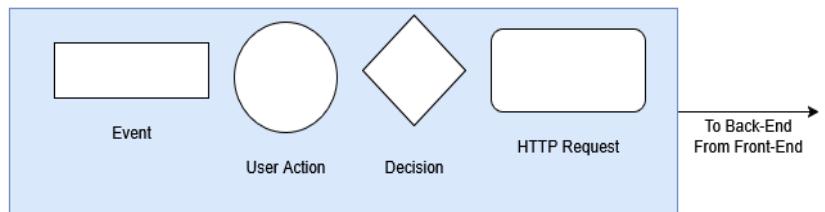
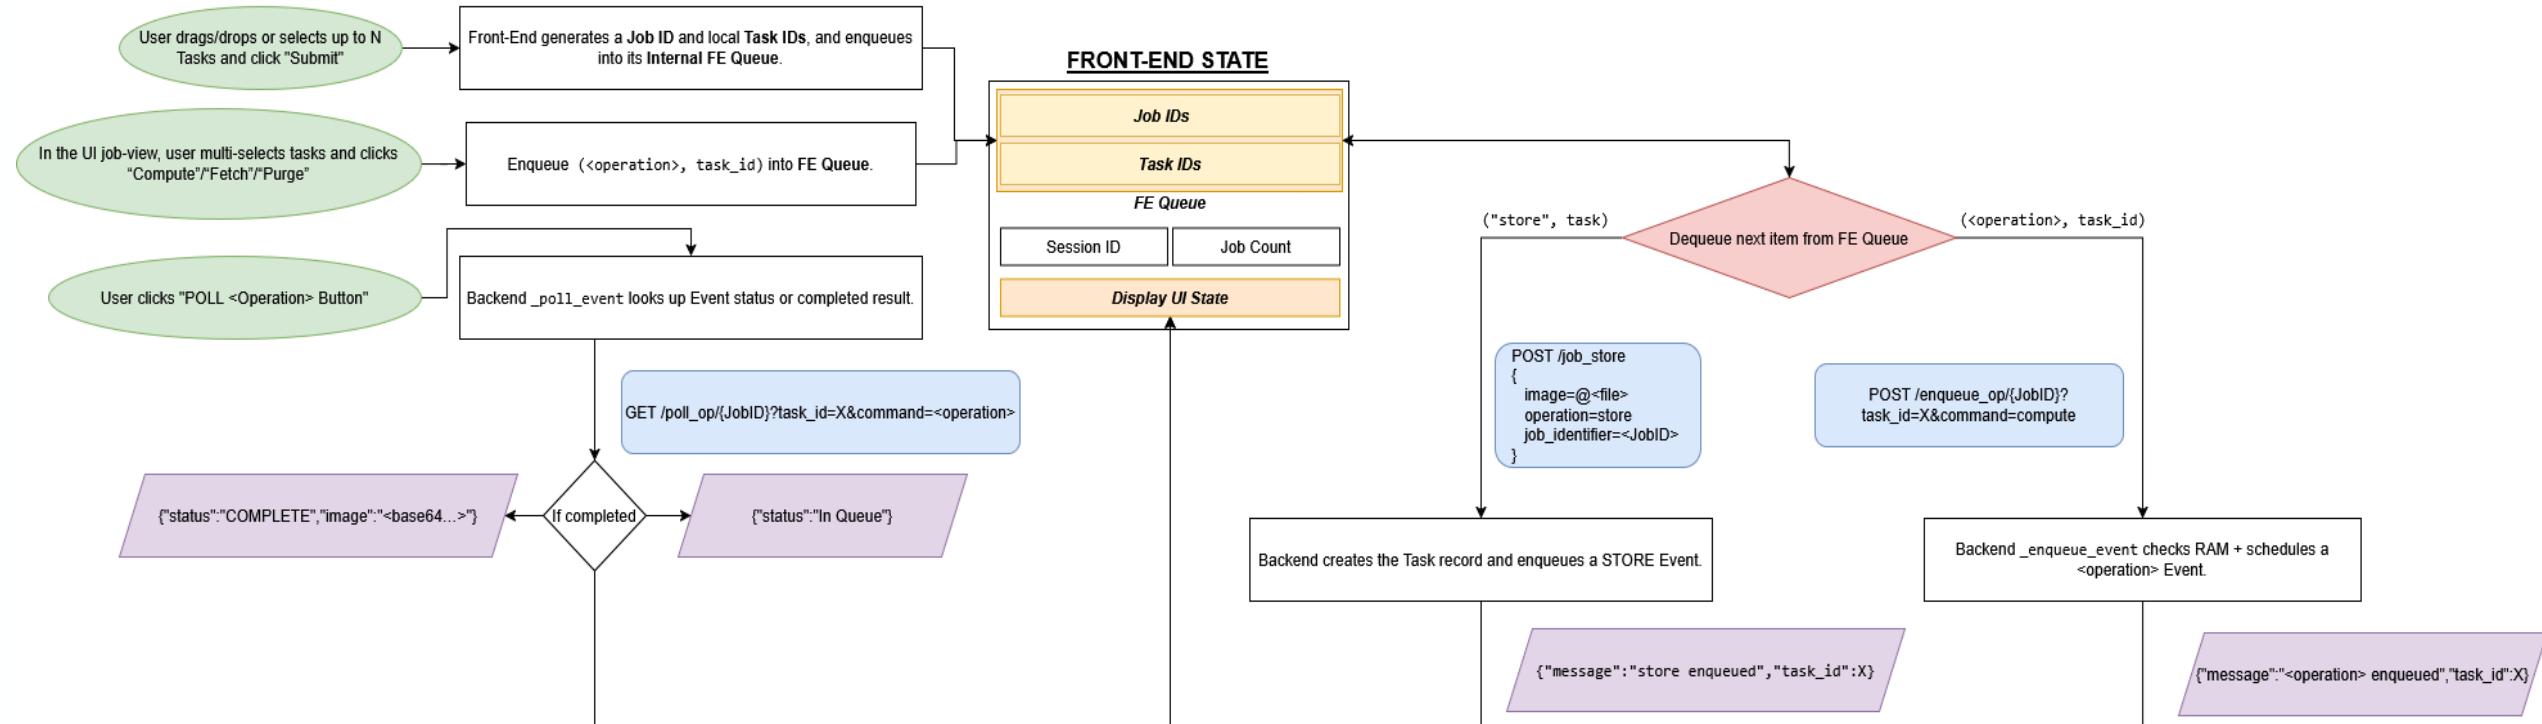
### Back-End Highlights:

- Built with FastAPI and AsyncIO for non-blocking operation handling
- Uses a multi-priority scheduler with RAM-aware enqueue policies
- Dedicated worker loop ensures COMPUTE and non-COMPUTE tasks are multiplexed

# User-Interface & Back-End Layer

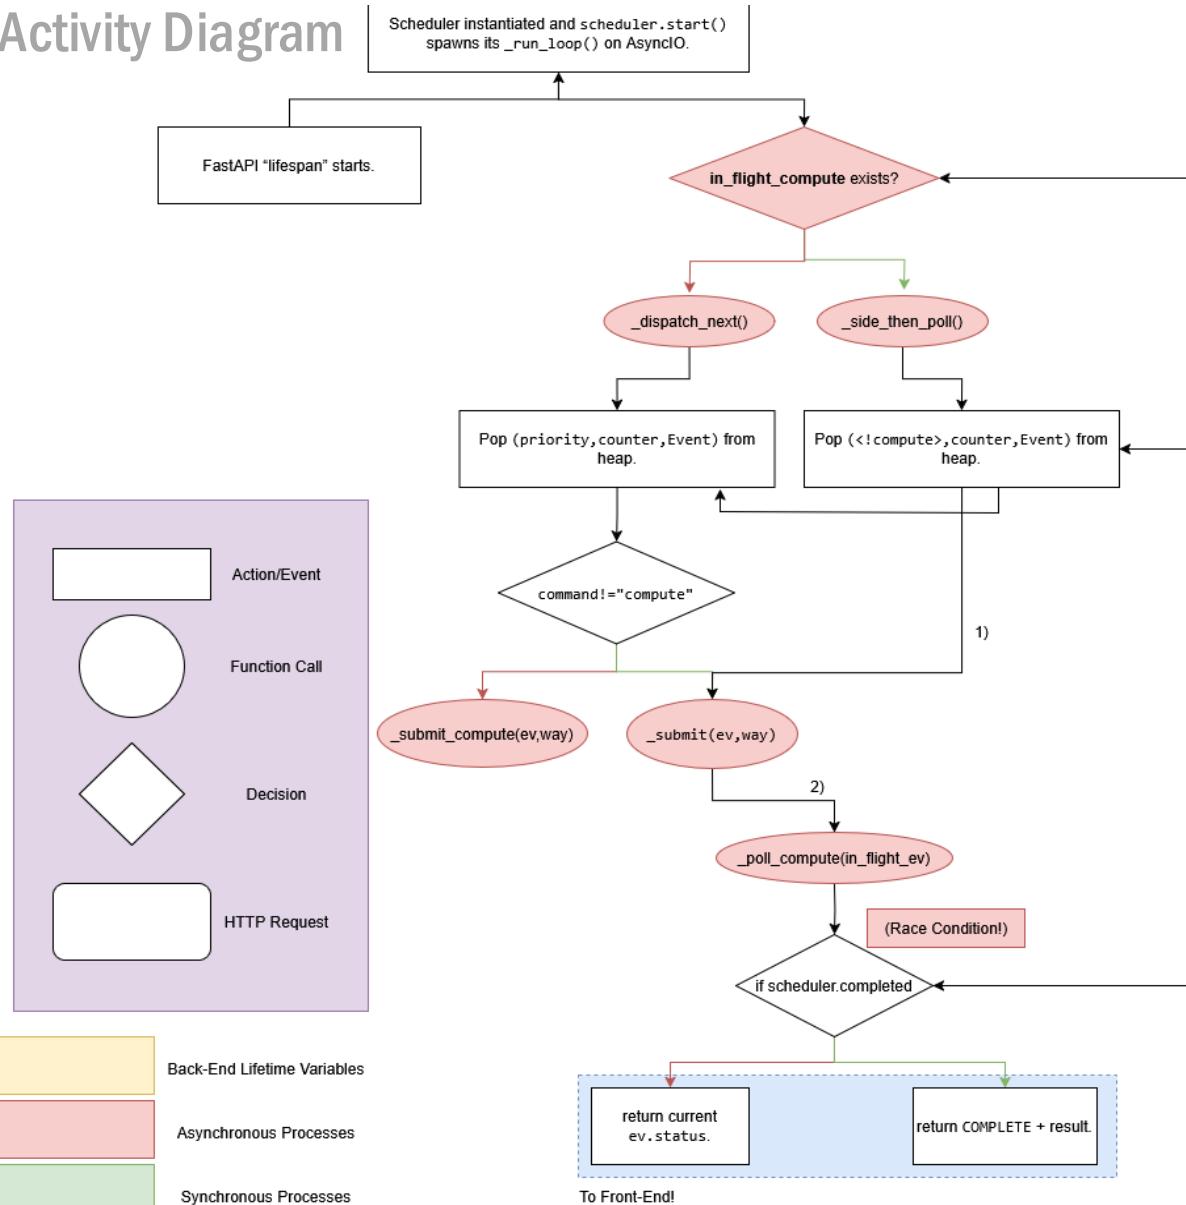
## Activity Diagram

### FRONT-END ACTIVITY DIAGRAM

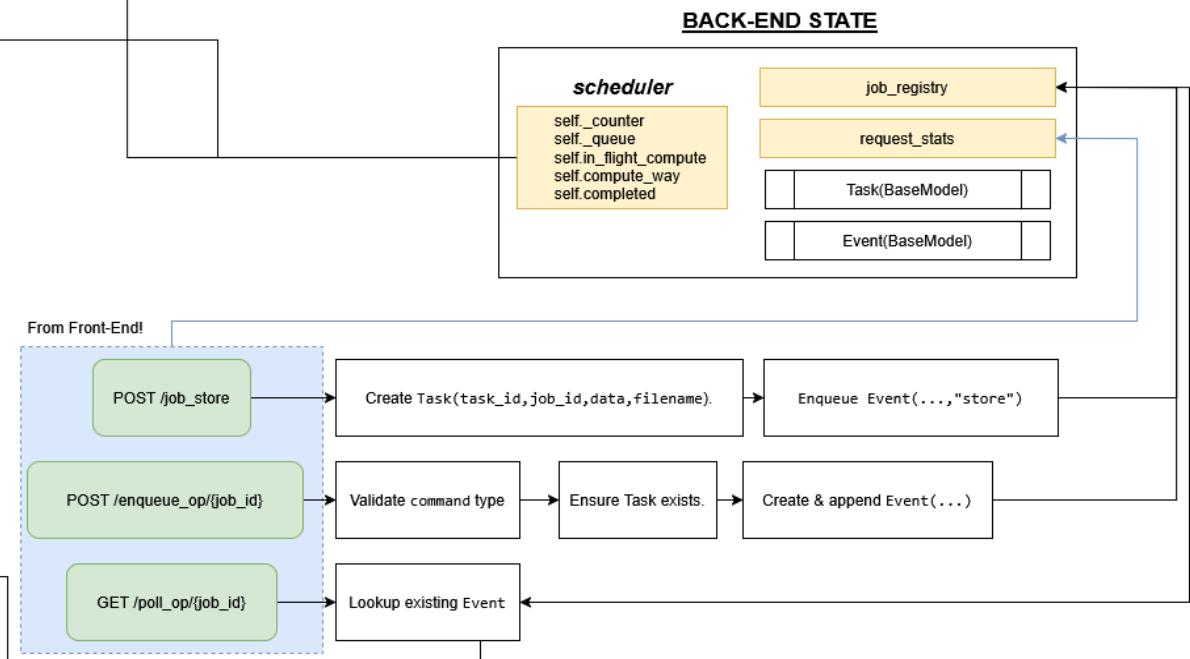


# User-Interface & Back-End Layer

## Activity Diagram



## BACK-END ACTIVITY DIAGRAM



# User-Interface & Back-End Layer

## Specifications

### Specification Relevance

- The Back-End must be able to guarantee response to any User-Interface request in under 50ms.
- The Back-End must prioritize STORE operations, and enable tracking & polling of in-flight COMPUTE operations.
- The Back-End must cancel transactions if total RAM usage exceeds 30% of available RAM
- The Back-End must be able to complete a 2KB transfer into the Network Node in under 100ms.

### Specification Testability

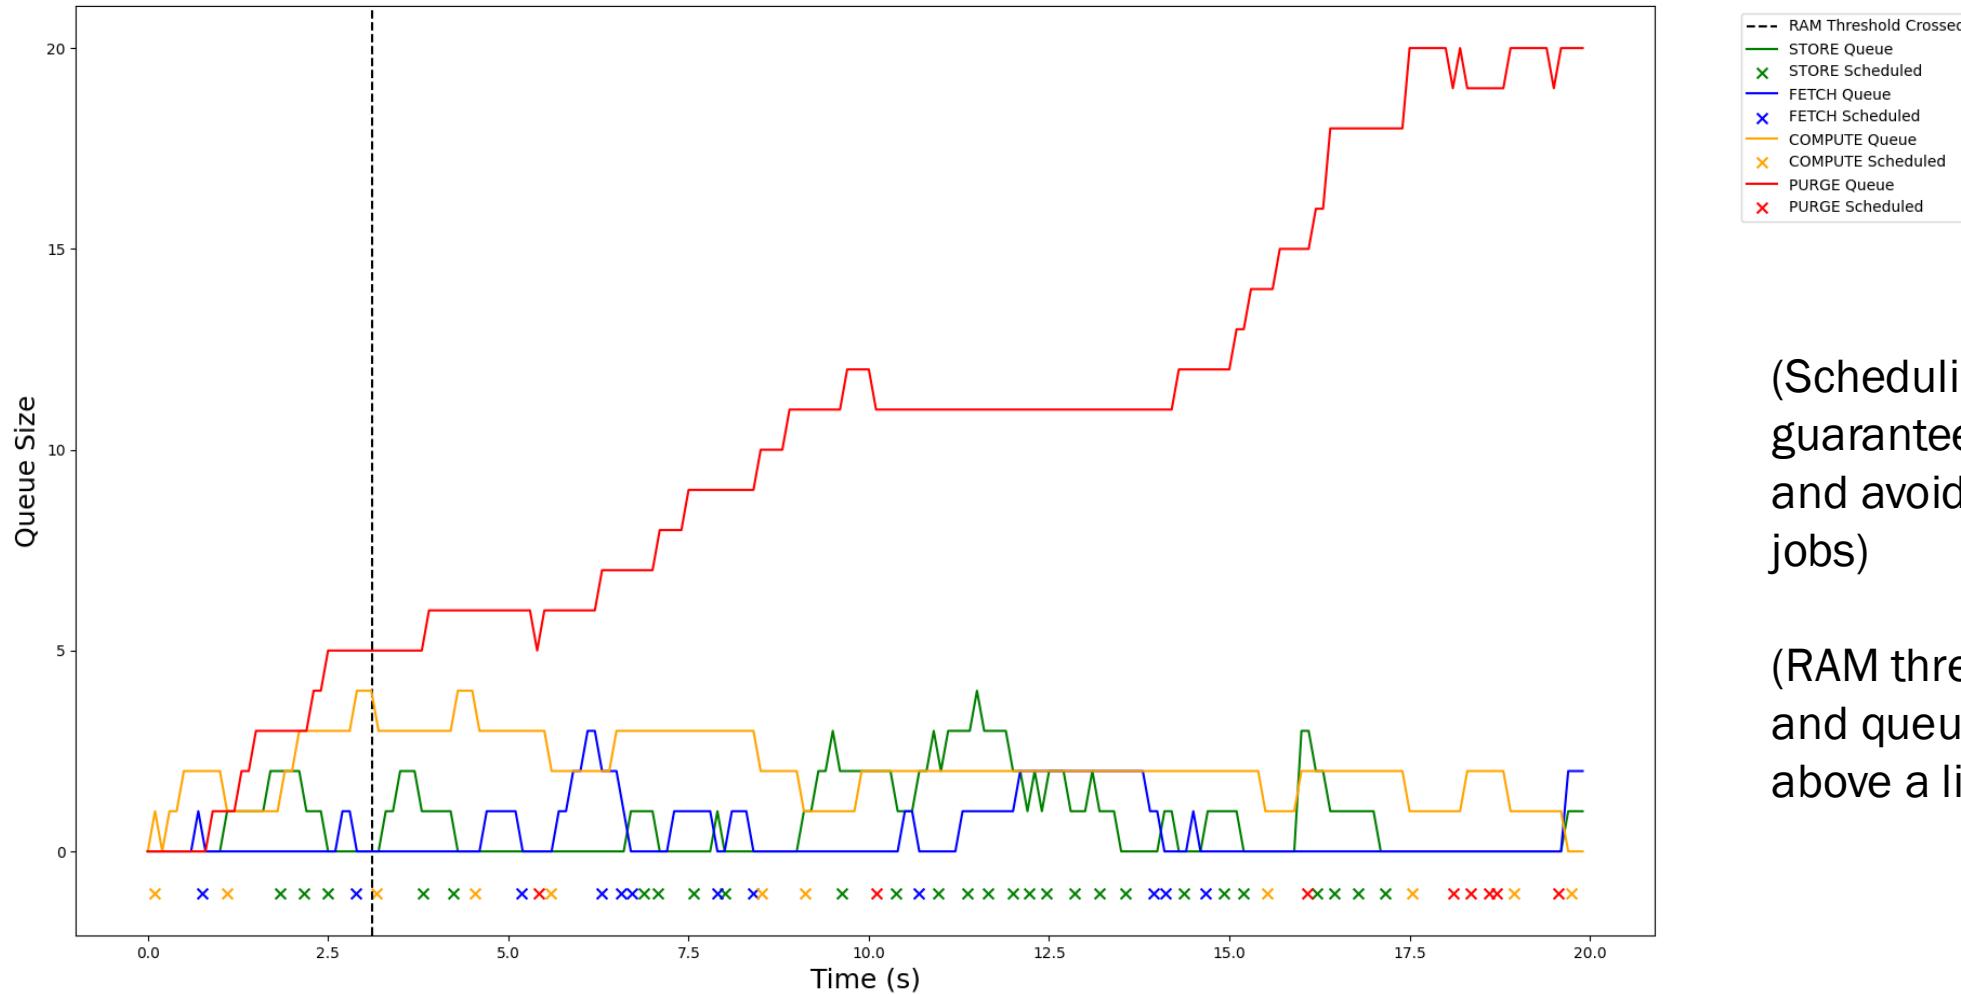
- Have a `app.middleware` endpoint, tracking lifetime of servicing all endpoints. Average out RTT throughout all requests.
- Stress Test the Scheduler layer; simulate randomized requests, delays and failures.
- Track RAM Usage in above Stress Test, and visualize peak queue-sizes.
- Timing Summaries across I/O heavy operations -- STORE and FETCH.

### Specification Motivation

- Ensures that the user interface remains responsive and avoids timeouts.
- Even under high system load, the scheduler ensures that STORE operations are not starved by longer-running commands. COMPUTE operations will be polled between trivial-ops.
- Prevents device failure!
- Guarantees minimal latency when synchronizing job/task updates between Backend and Nodes.

# User-Interface & Back-End Layer

## Specification Evidence





# User-Interface & Back-End Layer

## Specification Evidence



The diagram illustrates the interaction between the BoilerNet user interface and its back-end layer, showing how tasks are added, processed, and monitored.

**BoilerNet User Interface:**

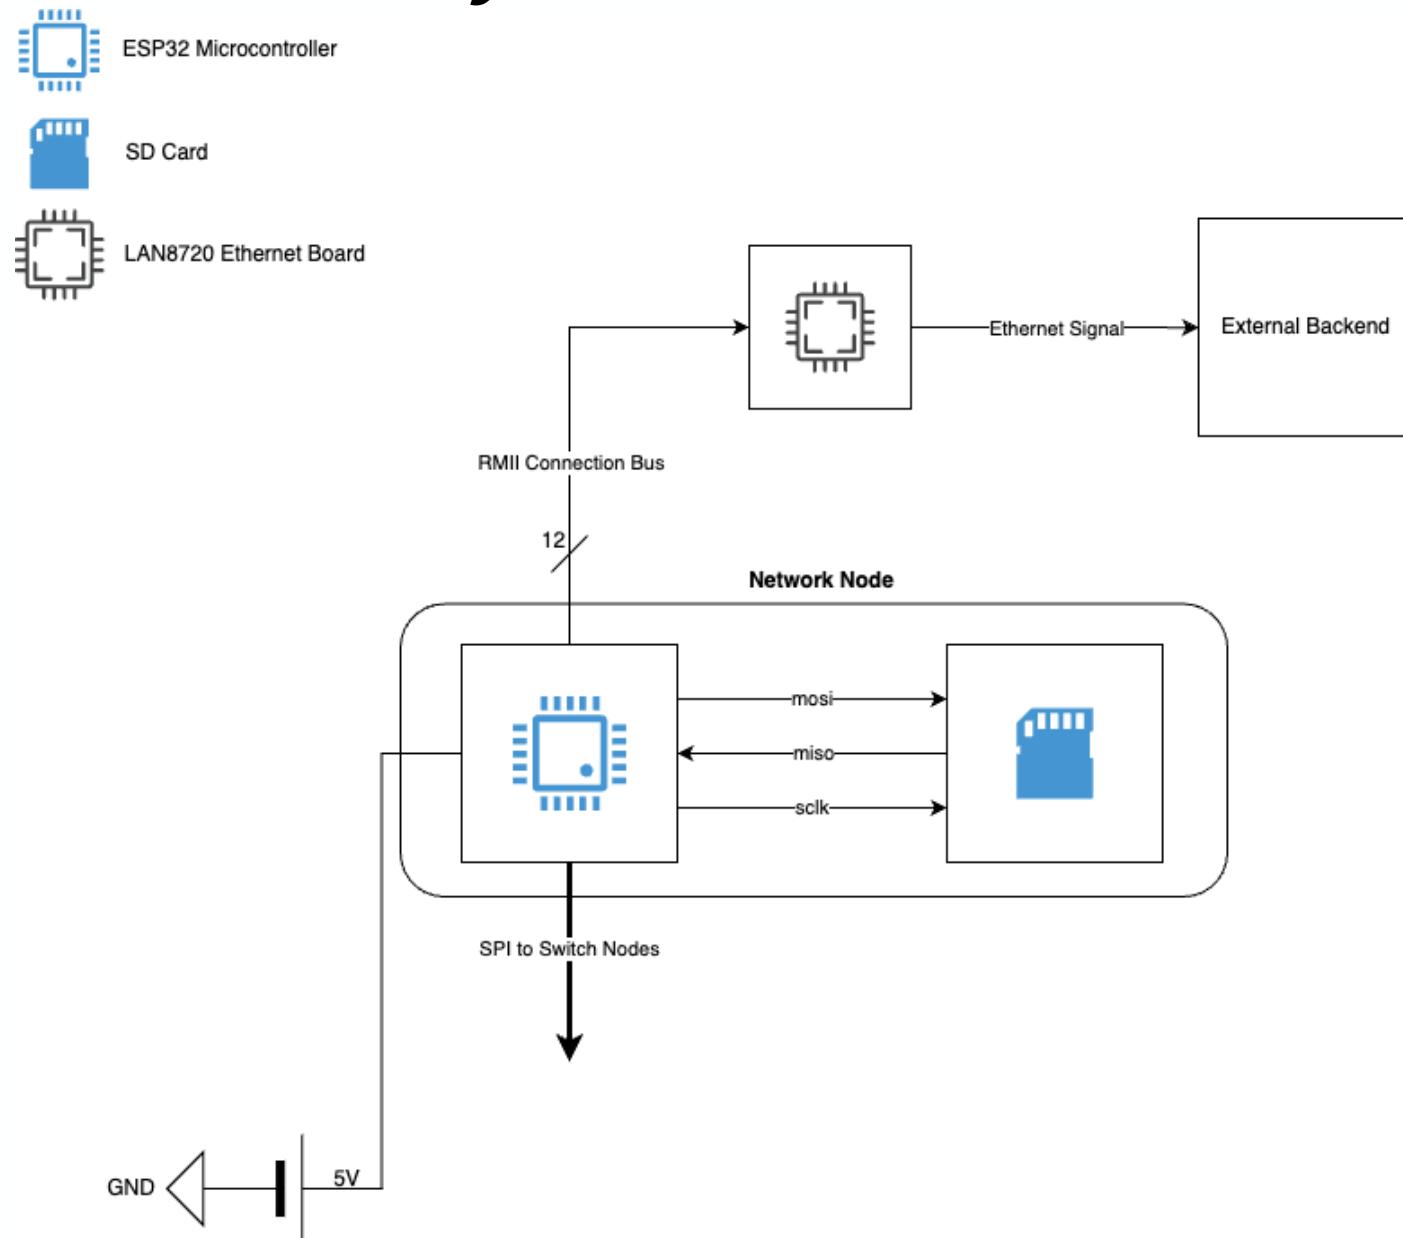
- Add New Job:** A screenshot showing three files being uploaded: `compute_working.png`, `write_network_timing.png`, and `read_network_timing.png`. A red box highlights the `Send jobs!>` button.
- Job Queue:** A screenshot showing a job named "Job 1". Under "Job 1 Tasks", two tasks are listed: `read_network_timing` and `write_network_timing`. An "Enqueue Operation" button is visible. Below the tasks, a message says "FETCH enqueued for task 1" and "FETCH enqueued for task 2".
- Poll Status:** A screenshot showing the status of Task 1. It indicates "Task 1 → COMPLETE" and displays a log of network activity. The log shows several `NEW_SOCKET` events, `SEND_PACKET` events, and `RECV_PACKET` events, with a total elapsed time of 167135 μs.
- Job Queue:** A screenshot showing the same "Job 1" and task list as the previous interface. The "Enqueue Operation" button is highlighted.
- Poll Status:** A screenshot showing the status of Task 1. It indicates "Task 1 → COMPLETE" and displays a log of network activity. The log shows several `NEW_SOCKET` events, `SEND_PACKET` events, and `RECV_PACKET` events, with a total elapsed time of 167135 μs.
- Poll Status:** A screenshot showing the status of Task 2. It indicates "Task 2 → COMPLETE" and displays a log of network activity. The log shows several `NEW_SOCKET` events, `SEND_PACKET` events, and `RECV_PACKET` events, with a total elapsed time of 15071 μs.

**Back-End Layer:**

- Task 1 → COMPLETE:** A detailed log of network activity for Task 1. It includes numerous `NEW_SOCKET`, `SEND_PACKET`, and `RECV_PACKET` events, along with timing information like "Elapsed: 167135 μs".
- Task 2 → COMPLETE:** A detailed log of network activity for Task 2. It includes numerous `NEW_SOCKET`, `SEND_PACKET`, and `RECV_PACKET` events, along with timing information like "Elapsed: 15071 μs".

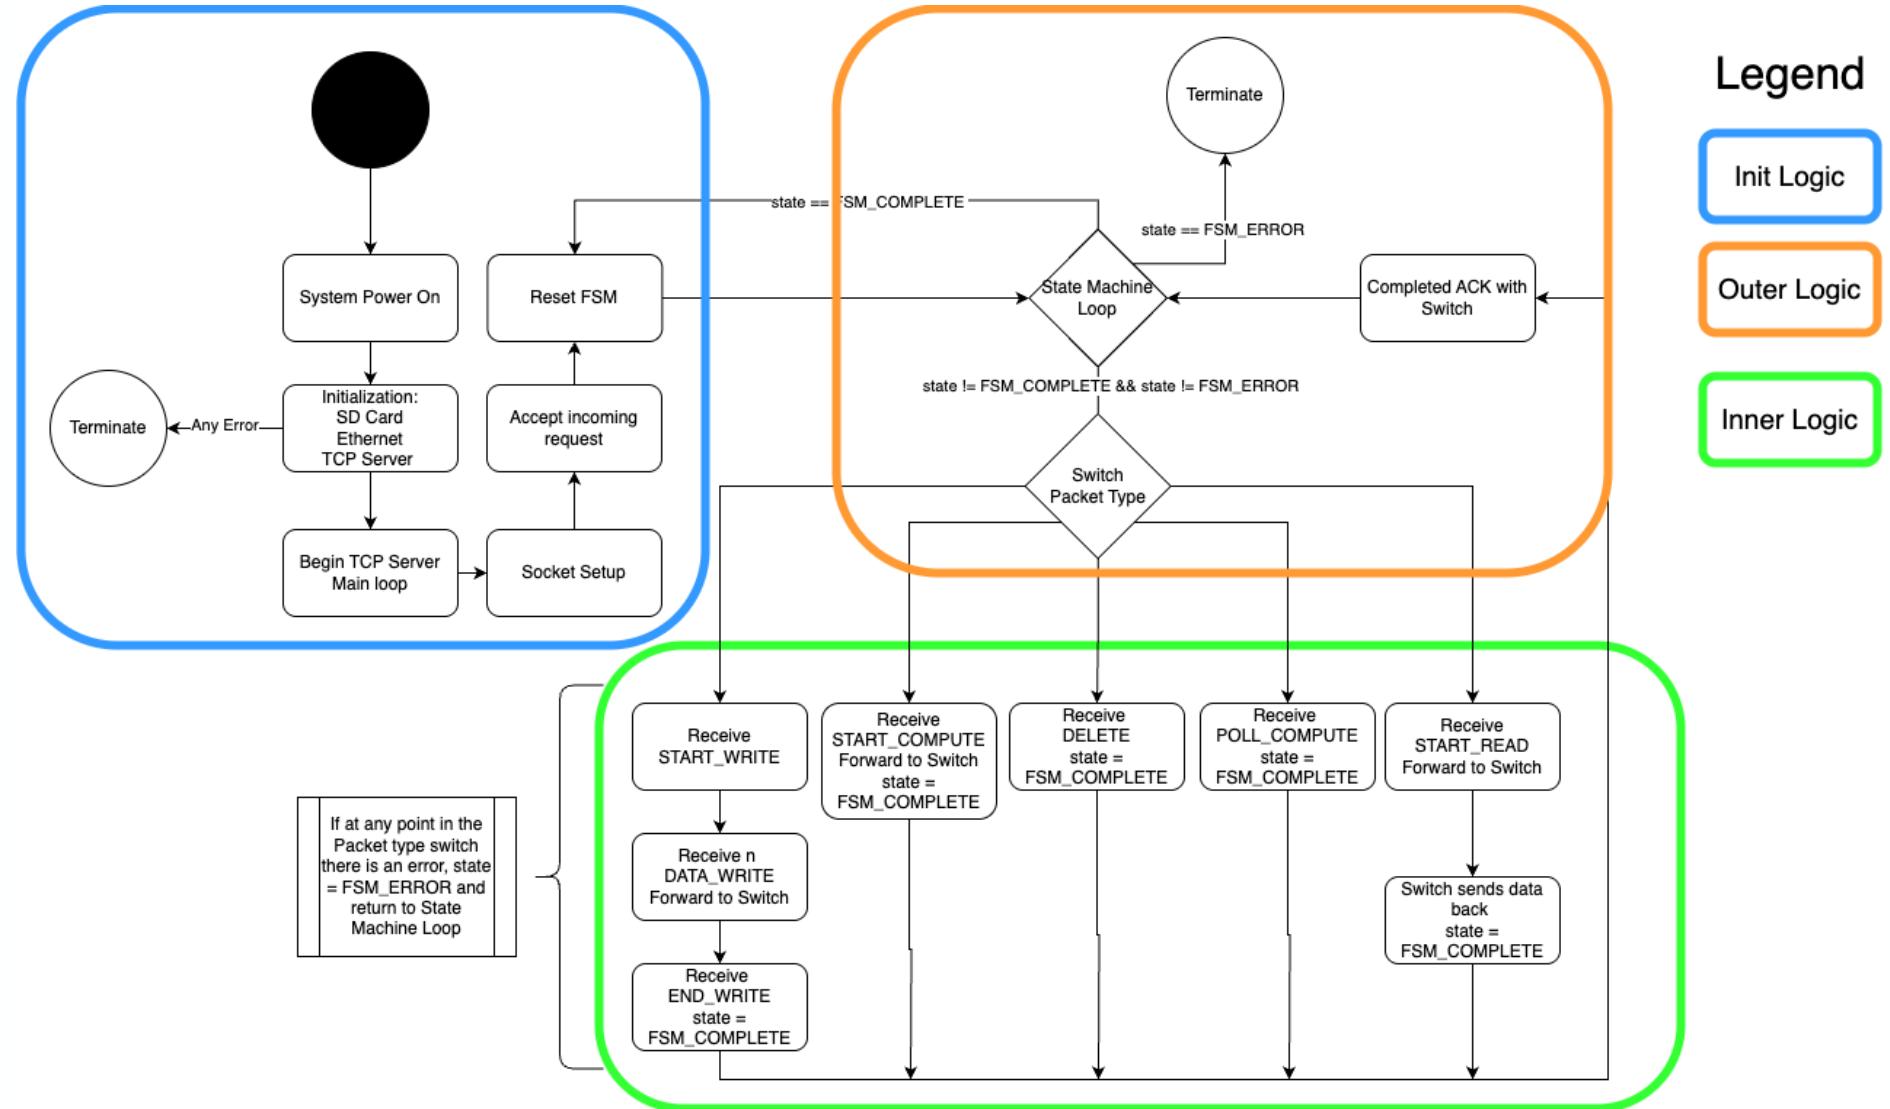
# *Network Coordination Layer*

Block Diagram



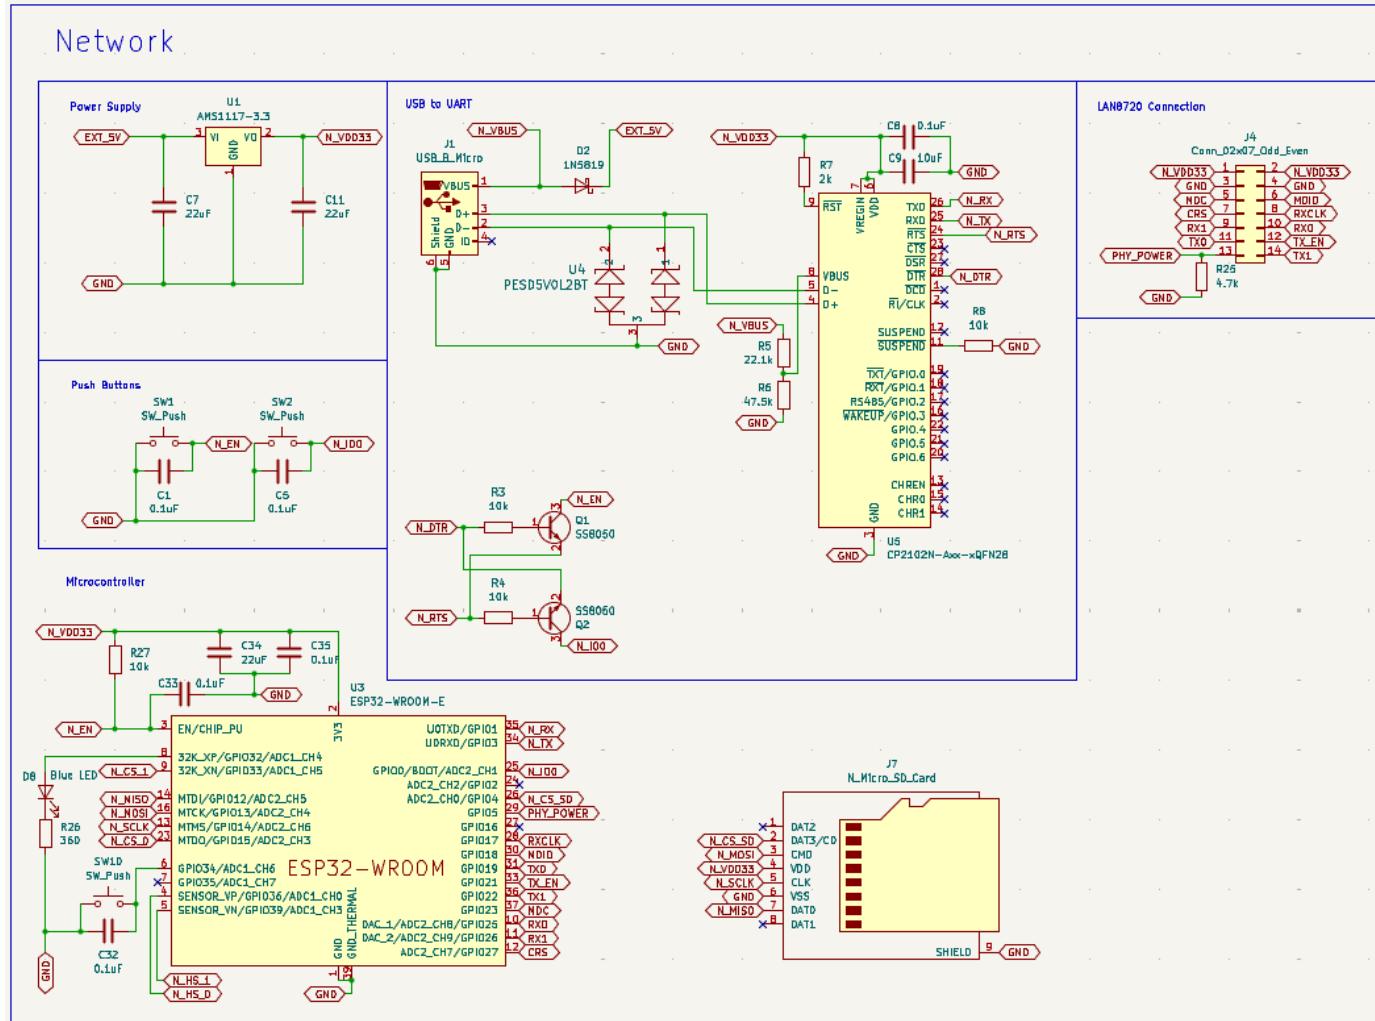
# Network Coordination Layer

Activity Diagram



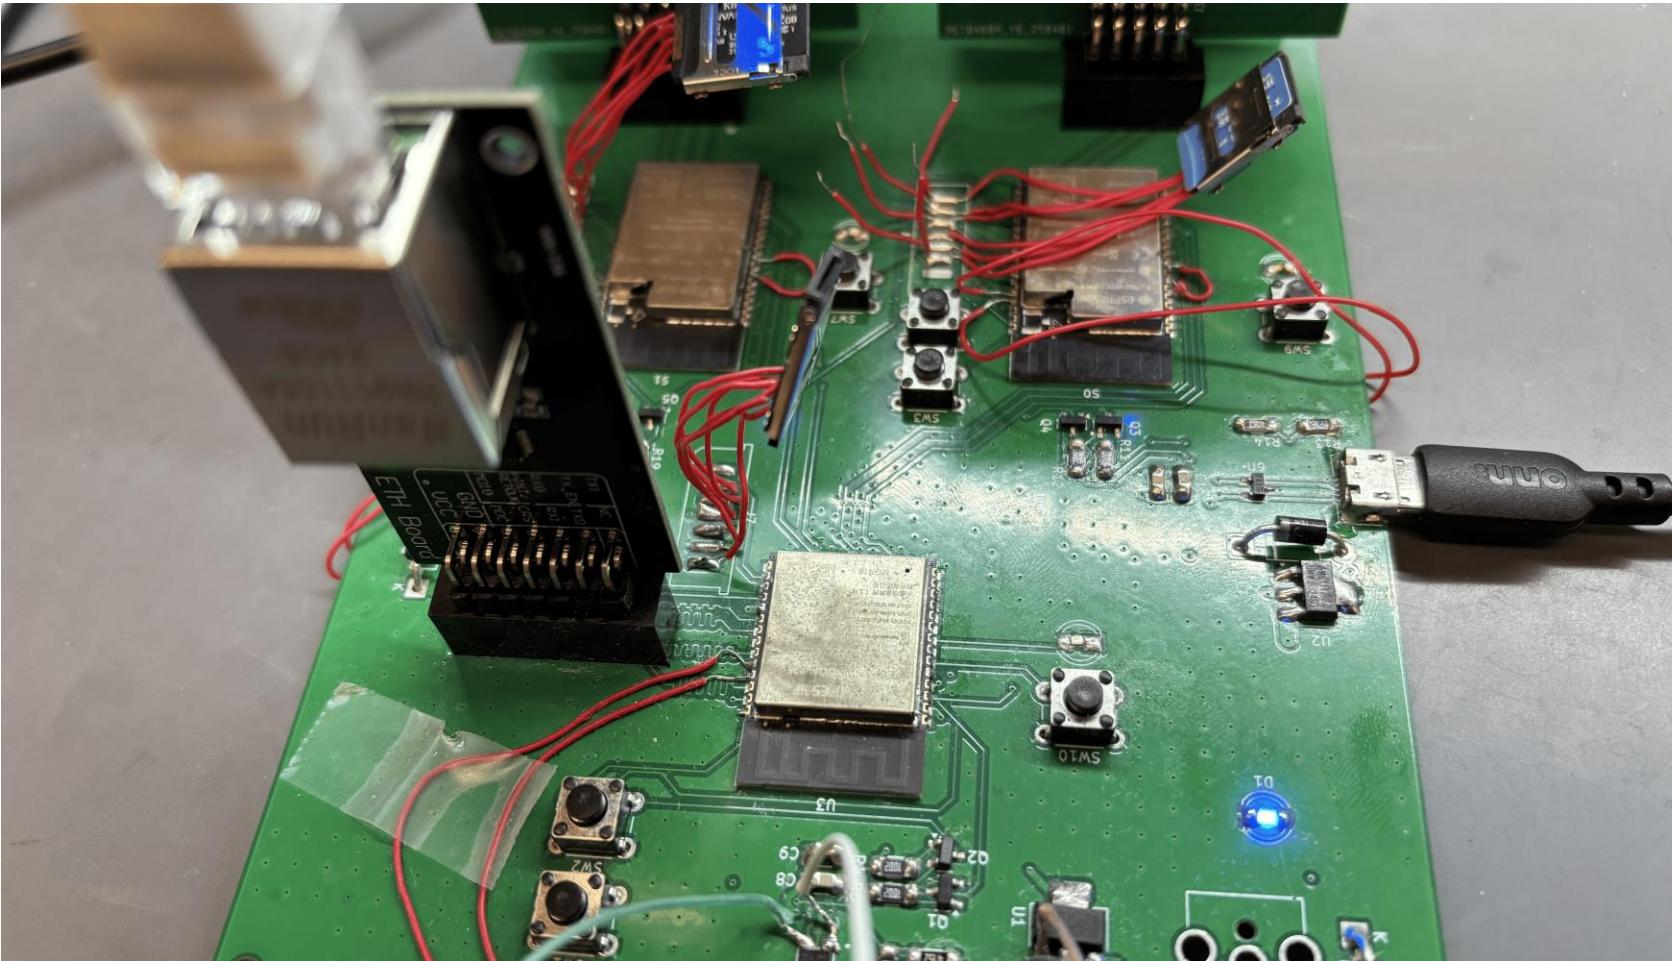
# Network Coordination Layer

## Schematic



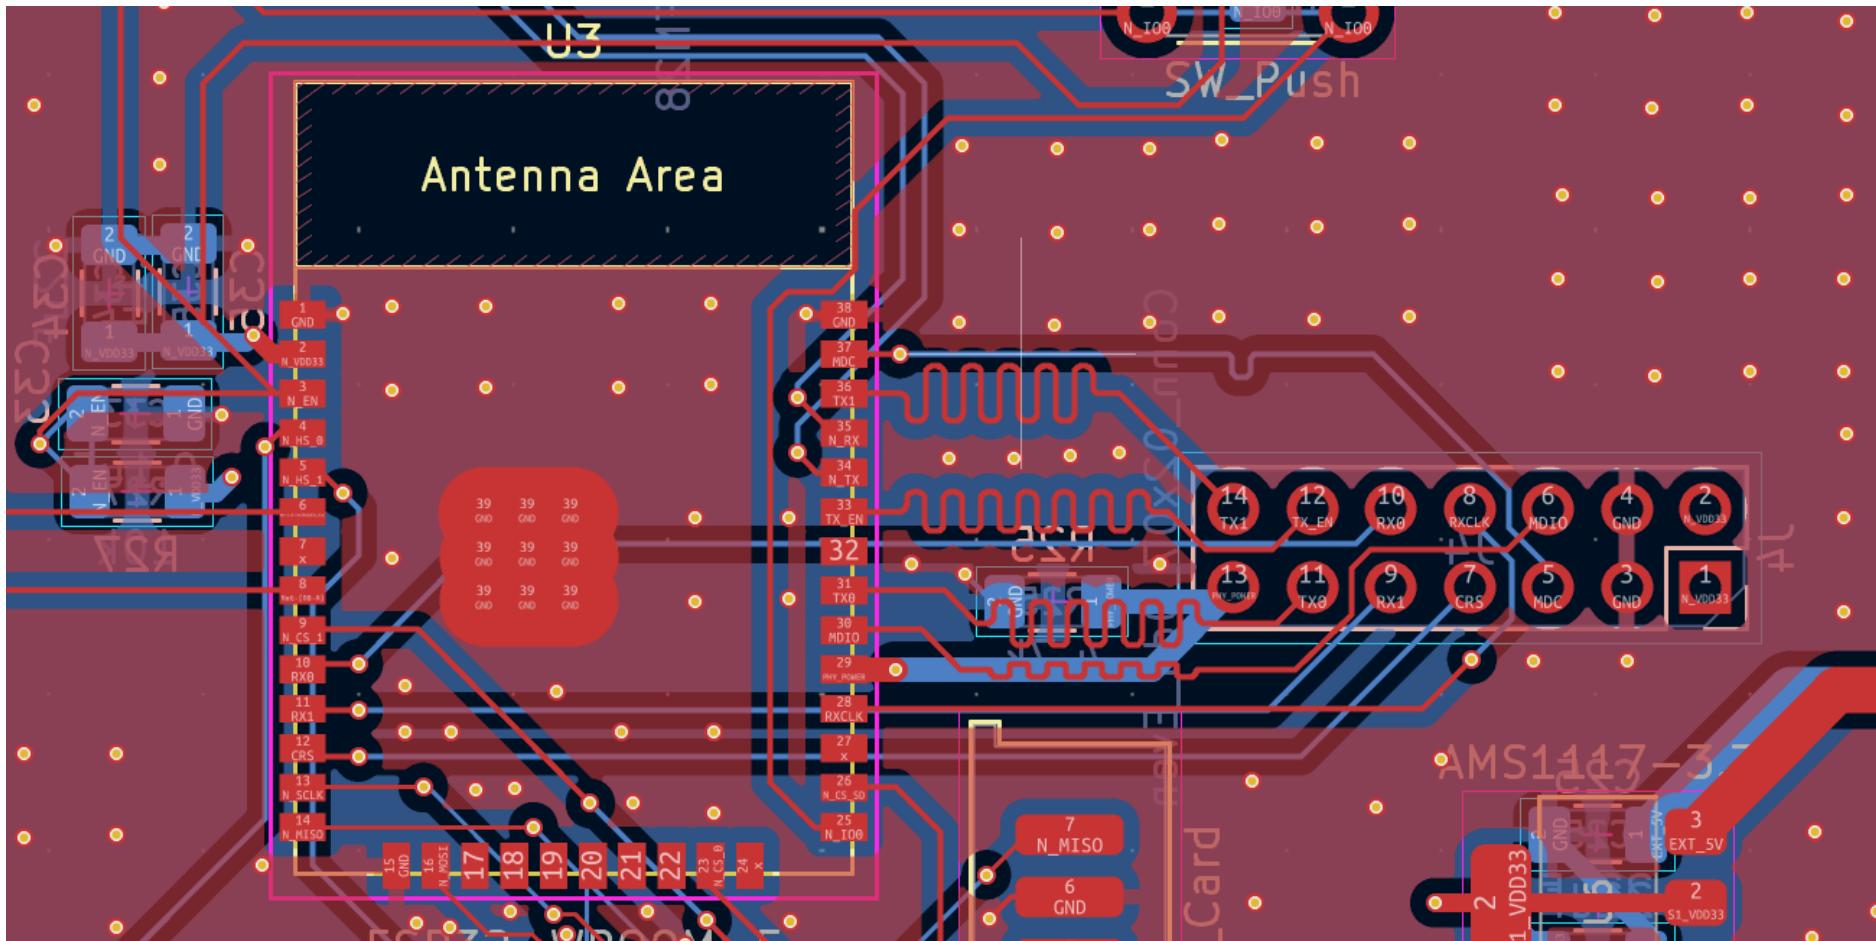
# *Network Coordination Layer*

PCB



# *Network Coordination Layer*

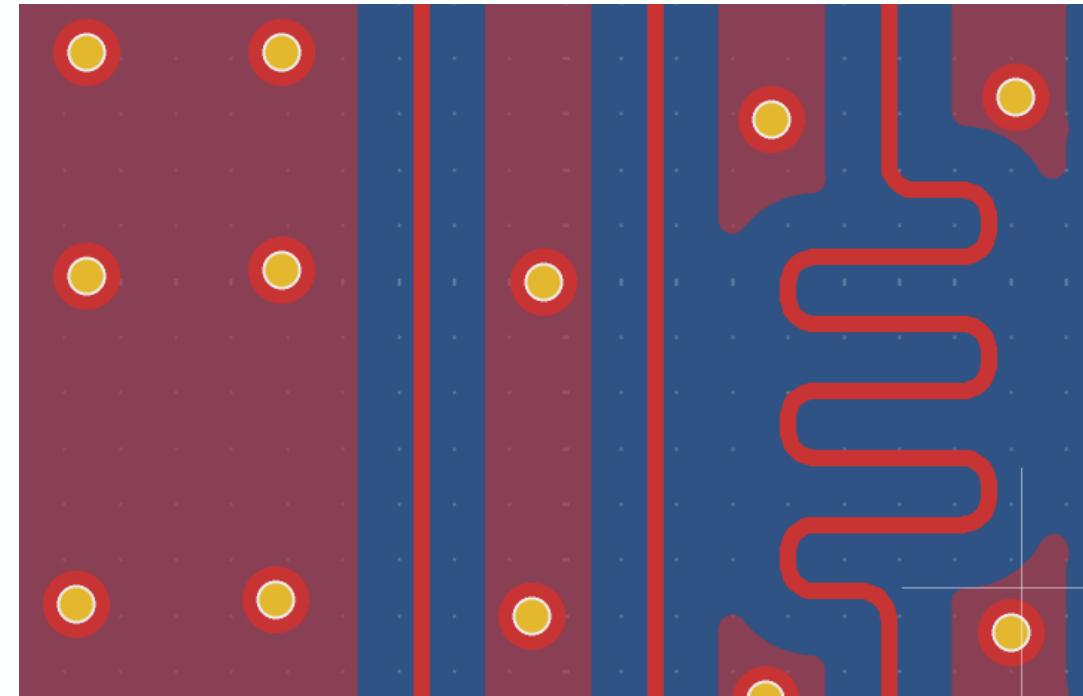
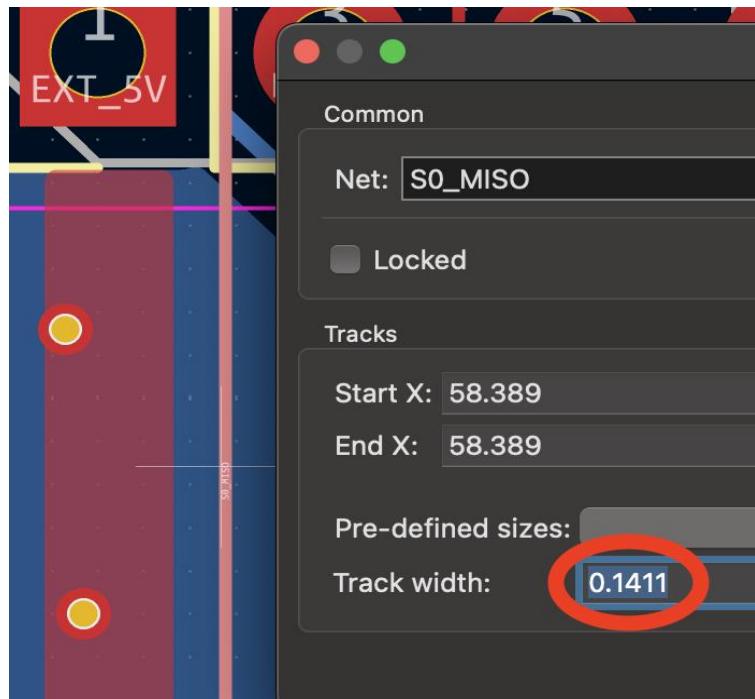
PCB



# *Network Coordination Layer*

PCB

- Length Matching
- Impedance Matching
- Via Stitching



# *Network Coordination Layer*

## Specification Relevance and Justification

- Average Response Time from the network for any Transaction should be less than 100ms.
- Average Response time for a full packet should be less than 2 seconds
- The Network Node must write to the SD Card at 20 Mhz

## Specification Testability

- We can test the Response time for a transaction from the backend side, timing the decision time and response time of the Network Node.
- We can time the full packet time for read and write operations by timing the system from the backend
- We can test the sd card write speed by using an oscilloscope.

## Specification Motivation

- The Backend should not be held up by exceedingly long processing times on the network node
- The write and read times should be independent of compute, and quick.
- Since BoilerNet is partly a storage solution, guaranteeing a minimum write speed to the sd card is needed.

# Specification 1 Measurement

Average Response Time from the Network Node for any Transaction should be less than 100ms.

```
I (86984) TCP_SOCKET: IN PACKET_TYPE_DATA_WRITE, data_len: 2030
I (86994) TCP_SOCKET: Sending buffer: 0xA1, 0x00 0x07 0xEE ... 0x00 0x00
I (87004) TCP_SOCKET: Received buffer: 0xA1, 0x00 0x07 0xEE ... 0x00 0x00
I (87004) TCP_SOCKET:      MATCH: SPI COMMAND: expected 0xA1, got 0xA1

Elapsed: 24929 µs
I (87014) TCP_SOCKET: IN PACKET_TYPE_DATA_WRITE, data_len: 2030
I (87334) TCP_SOCKET: Sending buffer: 0xA1, 0x00 0x07 0xEE ... 0x00 0x00
I (87344) TCP_SOCKET: Received buffer: 0xA1, 0x00 0x07 0xEE ... 0x00 0x00
I (87344) TCP_SOCKET:      MATCH: SPI COMMAND: expected 0xA1, got 0xA1

Elapsed: 333361 µs
I (87344) TCP_SOCKET: IN PACKET_TYPE_DATA_WRITE, data_len: 2030
I (87354) TCP_SOCKET: Sending buffer: 0xA1, 0x00 0x07 0xEE ... 0x00 0x00
I (87364) TCP_SOCKET: Received buffer: 0xA1, 0x00 0x07 0xEE ... 0x00 0x00
I (87364) TCP_SOCKET:      MATCH: SPI COMMAND: expected 0xA1, got 0xA1

Elapsed: 24901 µs
I (87374) TCP_SOCKET: !!!!~MID packet received~!!!!
I (87494) TCP_SOCKET: Sending buffer: 0xA2, 0x01 0x07 0xEE ... 0x00 0x00
I (87504) TCP_SOCKET: Received buffer: 0xA1, 0x00 0x07 0xEE ... 0x00 0x00
I (87504) TCP_SOCKET:      MATCH: SPI COMMAND: expected 0xA1, got 0xA1

Elapsed: 132359 µs
I (87514) TCP_SOCKET: IN PACKET_TYPE_DATA_WRITE, data_len: 2030
I (87514) TCP_SOCKET: Sending buffer: 0xA1, 0x00 0x07 0xEE ... 0x00 0x00
I (87524) TCP_SOCKET: Received buffer: 0xA2, 0x01 0x07 0xEE ... 0x00 0x00
I (87524) TCP_SOCKET:      MATCH: SPI COMMAND: expected 0xA2, got 0xA2

Elapsed: 20227 µs
I (87534) TCP_SOCKET: IN PACKET_TYPE_DATA_WRITE, data_len: 2030
I (87544) TCP_SOCKET: Sending buffer: 0xA1, 0x00 0x07 0xEE ... 0x00 0x00
I (87554) TCP_SOCKET: Received buffer: 0xA1, 0x00 0x07 0xEE ... 0x00 0x00
I (87554) TCP_SOCKET:      MATCH: SPI COMMAND: expected 0xA1, got 0xA1
```

```
I (8104) TCP_SOCKET: Received buffer: 0xB1, 0x07 0xFD 0x7D ... 0x0A 0x20
I (8104) TCP_SOCKET: DATA_READ packet: data length = 2045
I (8114) TCP_SOCKET: Forwarded DATA_READ packet, 2045 bytes
Elapsed: 26921 µs
I (8124) TCP_SOCKET:      Receiving Packet
I (8194) TCP_SOCKET: Sending buffer: 0xB1, 0x01 0x00 0x00 ... 0x00 0x00
I (8204) TCP_SOCKET: Received buffer: 0xB1, 0x07 0xFD 0x20 ... 0x28 0x6E
I (8204) TCP_SOCKET: DATA_READ packet: data length = 2045
I (8204) TCP_SOCKET: Forwarded DATA_READ packet, 2045 bytes
Elapsed: 89073 µs
I (8214) TCP_SOCKET:      Receiving Packet
I (8214) TCP_SOCKET: Sending buffer: 0xB1, 0x01 0x00 0x00 ... 0x00 0x00
I (8224) TCP_SOCKET: Received buffer: 0xB1, 0x07 0xFD 0x61 ... 0x66 0x6F
I (8224) TCP_SOCKET: DATA_READ packet: data length = 2045
I (8224) TCP_SOCKET: Forwarded DATA_READ packet, 2045 bytes
Elapsed: 26902 µs
I (8234) TCP_SOCKET:      Receiving Packet
I (8234) TCP_SOCKET: Sending buffer: 0xB1, 0x01 0x00 0x00 ... 0x00 0x00
I (8234) TCP_SOCKET: Received buffer: 0xB1, 0x07 0xFD 0x72 ... 0x6E 0x67
I (8234) TCP_SOCKET: DATA_READ packet: data length = 2045
I (8234) TCP_SOCKET: Forwarded DATA_READ packet, 2045 bytes
Elapsed: 89090 µs
I (8234) TCP_SOCKET:      Receiving Packet
I (8234) TCP_SOCKET: Sending buffer: 0xB1, 0x01 0x00 0x00 ... 0x00 0x00
I (8244) TCP_SOCKET: Received buffer: 0xB1, 0x07 0xFD 0x5F ... 0x0A 0x20
I (8244) TCP_SOCKET: DATA_READ packet: data length = 2045
I (8254) TCP_SOCKET: Forwarded DATA_READ packet, 2045 bytes
Elapsed: 26918 µs
I (8354) TCP_SOCKET:      Receiving Packet
I (8354) TCP_SOCKET: Sending buffer: 0xB1, 0x01 0x00 0x00 ... 0x00 0x00
I (8434) TCP_SOCKET: Received buffer: 0xB1, 0x07 0xFD 0x20 ... 0x61 0x22
I (8434) TCP_SOCKET: DATA_READ packet: data length = 2045
I (8444) TCP_SOCKET: Forwarded DATA_READ packet, 2045 bytes
```

# *Specification 2 Measurement*

Average Response Time for a full packet should be less than 2 s.

```
● (base) root@araviki:/home/araviki/workbench/boilernet/test# python3 test_routines.py 8080 192.168.0.50 chiahua.jpg read user02
OPERATION: read | USER_ID: user02
Connecting over Socket
Opened Socket!
READ complete, checksum=00AD9339

==== Timing Summary ====
Total: 2994.49 ms
Avg pkt time: 7.36 ms
    ACK_B0: 18.77 ms
    DATA READ: 7.10 ms

● (base) root@araviki:/home/araviki/workbench/boilernet/test# python3 test_routines.py 8080 192.168.0.50 chiahua.jpg write user02
OPERATION: write | USER_ID: user02
Connecting over Socket
Opened Socket!
STORE on FILENAME: ./input/chiahua.jpg | ext: .jpg
Total Bytes: 89604 | Packets: 45
Sent END packet (write) with checksum 0x00ad9339
WRITE complete

==== Timing Summary ====
Total: 4486.31 ms
Avg pkt time: 1478.80 ms
    ACK_A0: 12.25 ms
    ACK_A2: 2749.31 ms
    ACK A3: 1674.84 ms
```

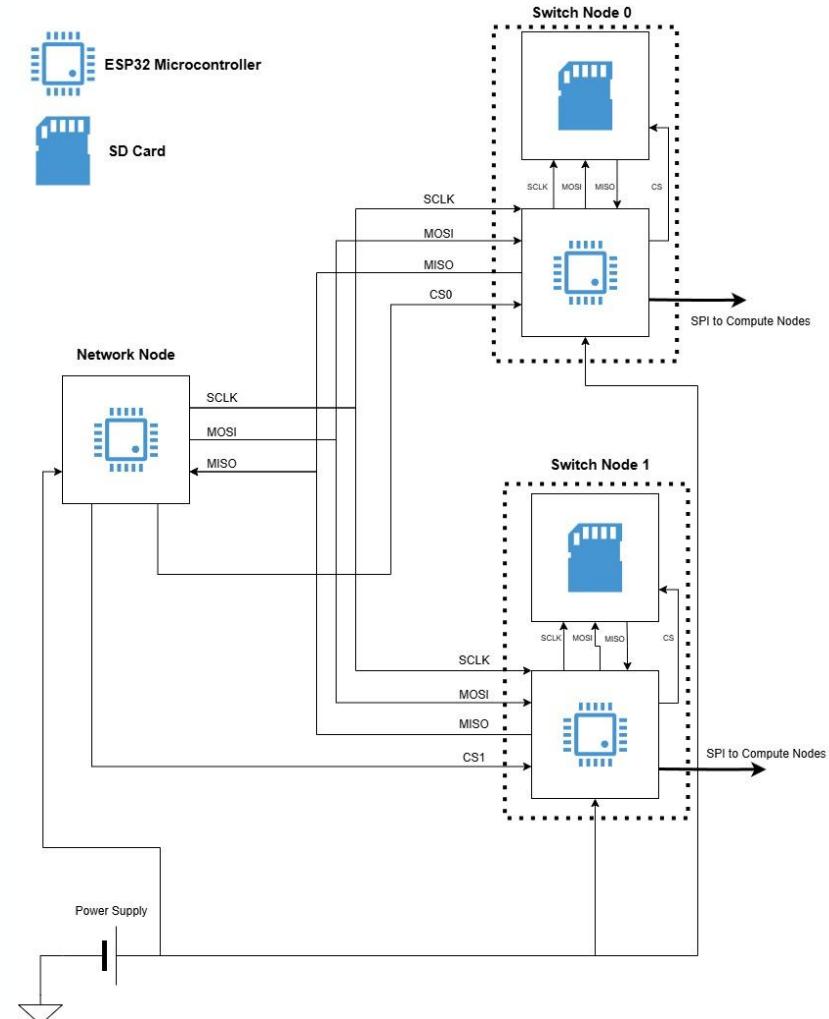
# Specification 3 Measurement

The Network Node must write to the SD Card at 20 Mhz



# *Internal Switch Routing Layer*

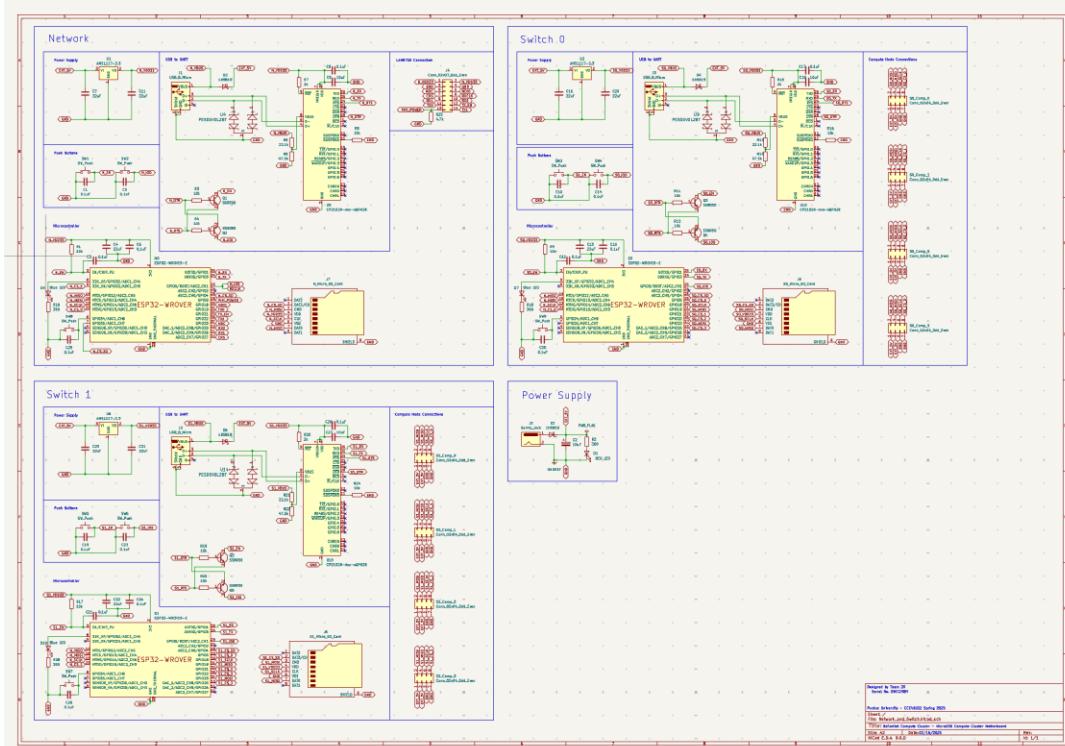
## Block Diagram



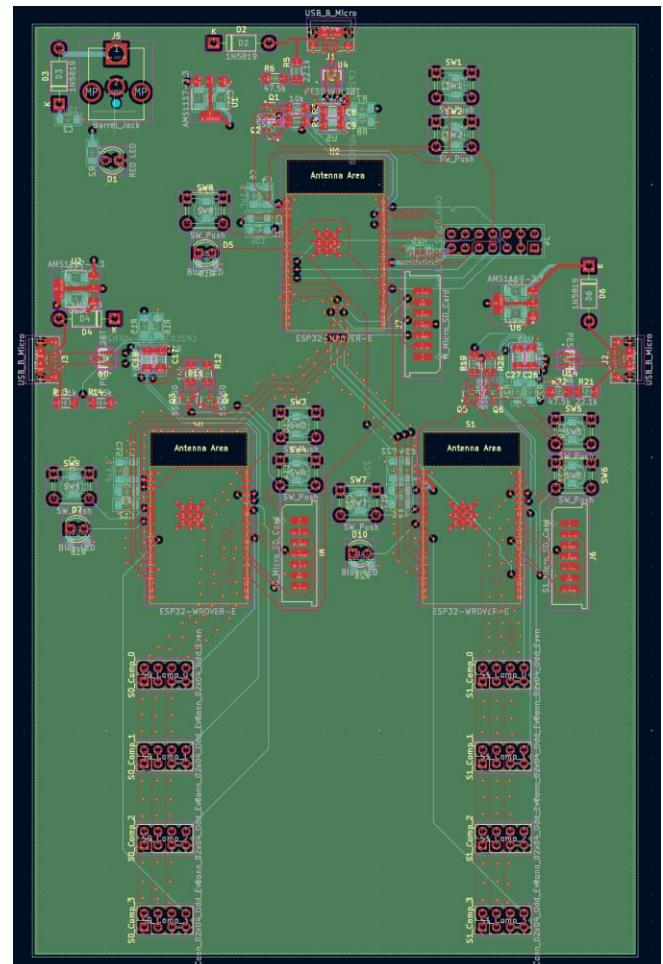
# *Internal Switch Routing Layer*

Schematic

Network

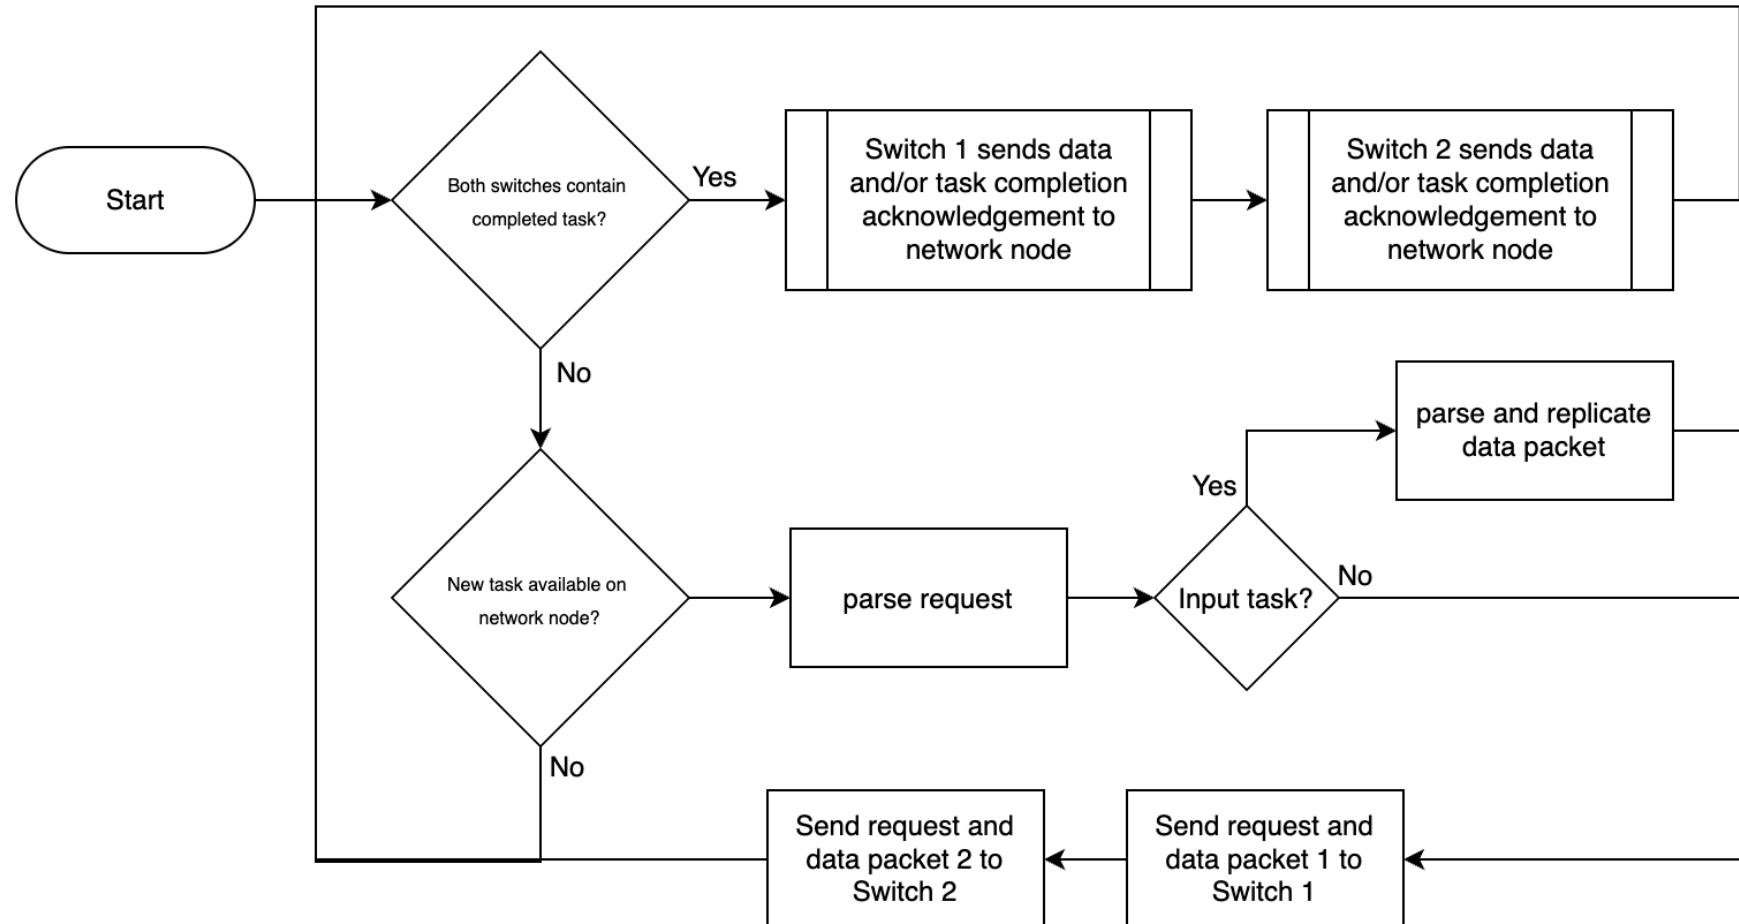


Switch 0



# *Internal Switch Routing Layer*

## Activity Diagram



# *Internal Switch Routing Layer*

## Specification Relevance and Justification

- **Relevance:** Groups the compute nodes using switches to allow for fewer slave devices per communication master and allows data placement over multiple storage devices.
- **Justification:** Allowing the cluster to store data across different storage mediums on different nodes allows for much greater file storage. Utilizing two different switch nodes also lets the cluster scale more effectively and minimize congestion on any one communication bus.

## Specification Testability

- Speed of subsystem and subsystem parts can be tested by measuring the time it takes between a command and data being received by the network node and the second switch node ready to send the data to its compute nodes.
- Check to see if error packets are being received successfully
- Check for data persistence
- Check for critical timing delay with a similar method to subsystem speed measurement

## Specification Motivation

- The Internal Switch Routing Layer was devised in order to add a hierarchical level to the compute cluster for boosting cluster scalability and for introducing data storage striping across multiple storage mediums.

# *Specifications*

Blue – MET, Yellow – PARTIALLY MET, Red - NOT MET

- Individual packets can be received and sent at greater than 50KBps over SPI.
- The Critical Path Delay in a Switch -- from receiving a packet, to preparing for the next one -- must be under 200ms.
- The Switch Nodes must be able to communicate errors to both the Network Node and Compute Nodes.
- The Switch Nodes must maintain data when there is a power loss.

# *Specification Measurement*

Individual packets can be received and sent at greater than 50KBps over

```
[recv] DATA_READ_LEN: 2045 bytes -> 0.006s
[recv] DATA_READ_LEN: 2045 bytes -> 0.005s
[recv] DATA_READ_LEN: 1531 bytes -> 0.005s
[recv] END_READ_CHECKSUM: 0x00 0x83 0xCE 0xC0
READ complete
    Total bytes: 66971
    Duration: 2.518s
    Throughput (including I/O + Control Logic): 26600.51 bytes/s
    Checksum: 0083CEC0

    === Timing Summary ===
    Total: 2548.58 ms
    Avg pkt time: 9.15 ms
        ACK_B0: 22.13 ms
        DATA_READ: 8.76 ms
```

26.6 KB/s

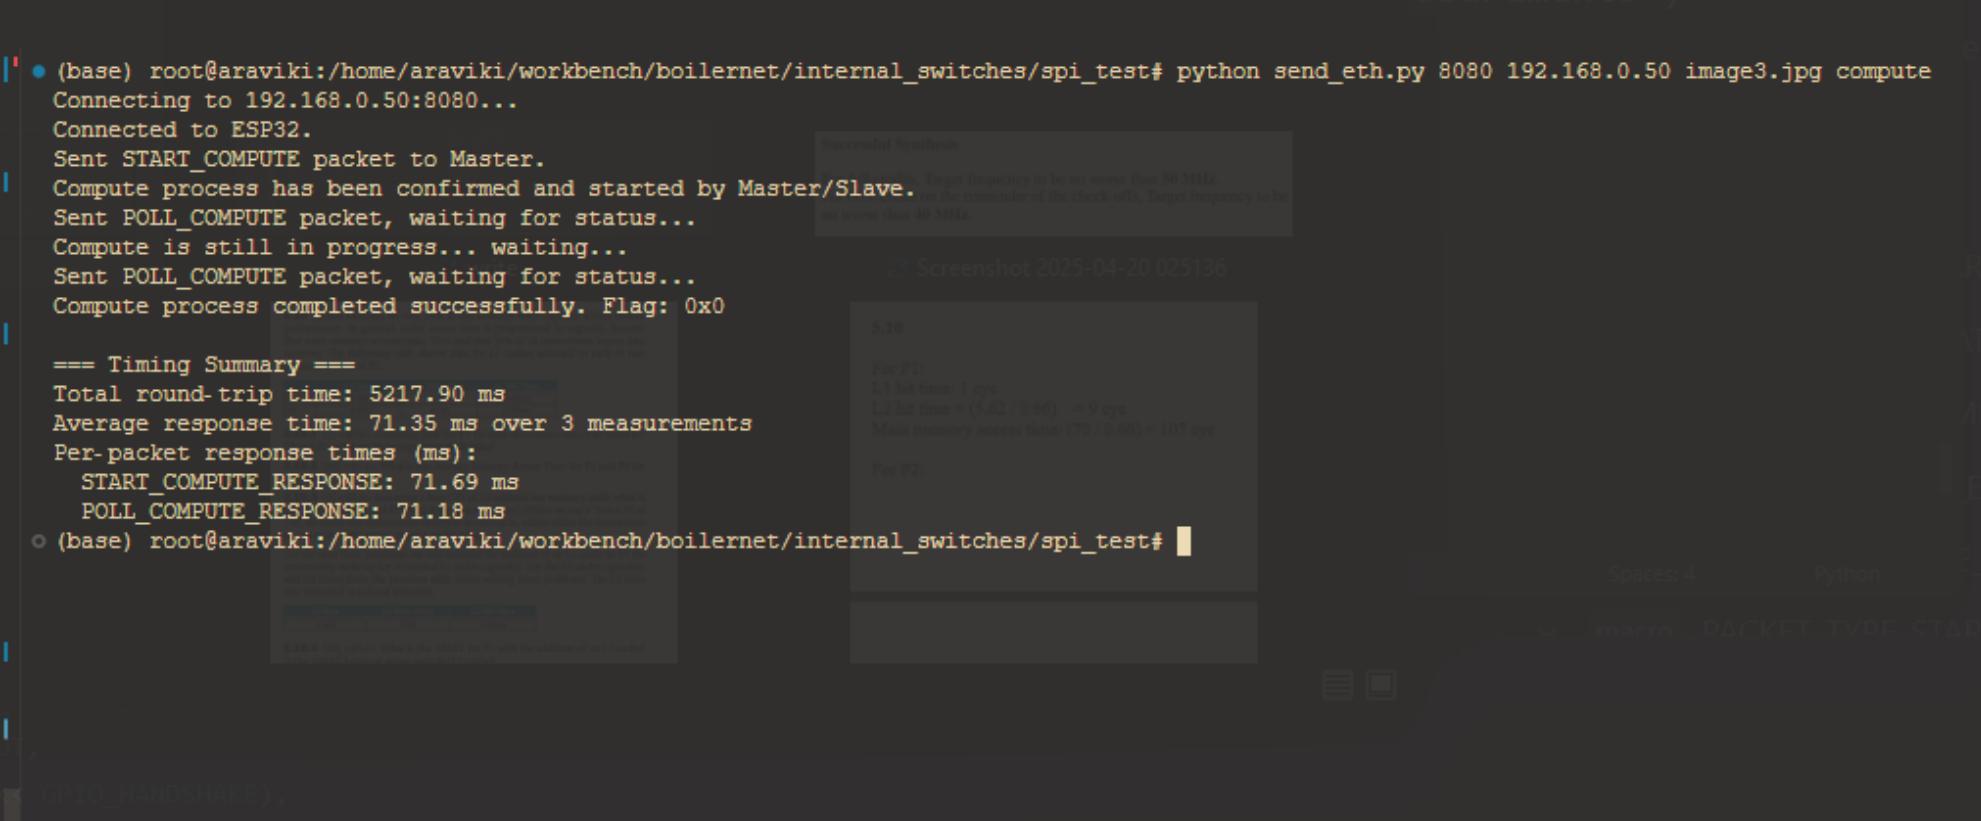
```
[recv] DATA_READ_LEN: 2050 bytes -> 0.000s
[recv] DATA_READ_LEN: 2030 bytes -> 0.000s
[recv] DATA_READ_LEN: 2011 bytes -> 0.000s
Sent END packet (write) with checksum 0x0083cec0
WRITE complete
    Duration: 3.623s
    Throughput: 18487.04 bytes/s

    === Timing Summary ===
    Total: 3623.03 ms
    Avg pkt time: 100.59 ms
        ACK_A0: 100.51 ms
        DATA_WRITE: 0.01 ms
        ACK_A2: 2213.27 ms
        ACK_A3: 1307.28 ms
```

18.5 KB/s

# Specification Measurement

The Critical Path Delay in a Switch – from receiving a packet, to preparing for the next one – must be under 200ms.



```
● (base) root@araviki:/home/araviki/workbench/boilernet/internal_switches/spi_test# python send_eth.py 8080 192.168.0.50 image3.jpg compute
Connecting to 192.168.0.50:8080...
Connected to ESP32.
Sent START_COMPUTE packet to Master.
Compute process has been confirmed and started by Master/Slave.
Sent POLL_COMPUTE packet, waiting for status...
Compute is still in progress... waiting...
Sent POLL_COMPUTE packet, waiting for status...
Compute process completed successfully. Flag: 0x0

==== Timing Summary ====
Total round-trip time: 5217.90 ms
Average response time: 71.35 ms over 3 measurements
Per-packet response times (ms):
  START_COMPUTE_RESPONSE: 71.69 ms
  POLL_COMPUTE_RESPONSE: 71.18 ms
● (base) root@araviki:/home/araviki/workbench/boilernet/internal_switches/spi_test#
```

Successful Synthesis

Screenshot 2025-04-20 025136

For P1:  
L1 hit time: 1 cyc  
L2 hit time =  $(8.62 / 0.66) = 9$  cyc  
Main memory access time:  $(70 / 0.66) = 107$  cyc

For P2:

Spaces: 4    Python

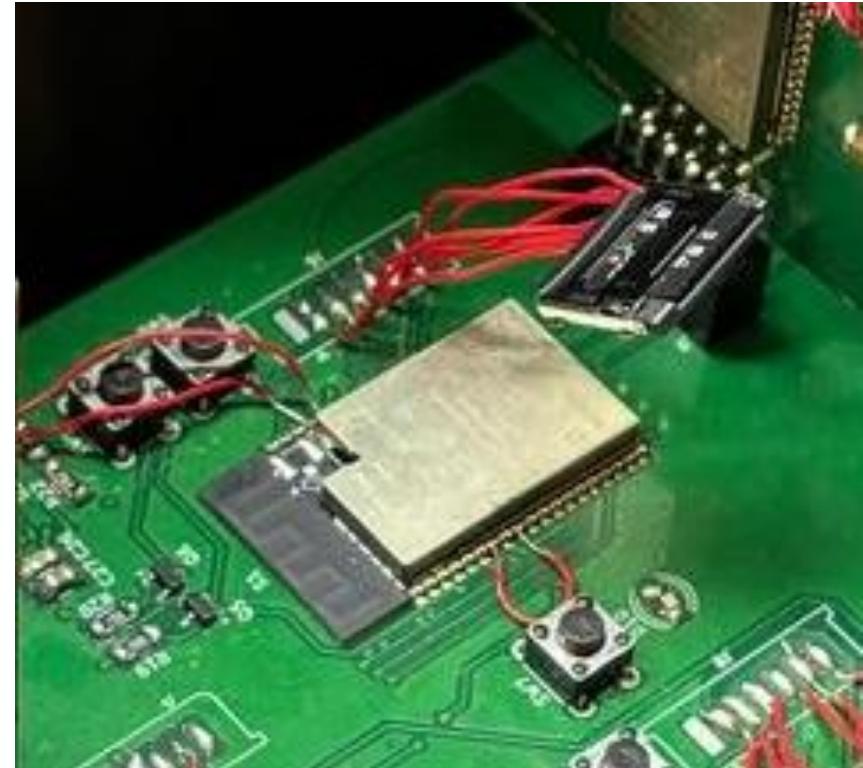
macro PACKET\_TVDF STAR

GPIO\_HANDSHAKE),

# *Specification Measurement*

```
(base) root@araviki:/home/araviki/workbench/boilernet/internal_switches/s  
age3.jpg write  
Connecting to 192.168.0.50:8080...  
Connected to ESP32.  
Image loaded: 25x25, total 9155 bytes  
num_full_packets 4  
remainder 1035  
total_packets 5  
mid_packet_index 2  
Sent START packet for write.  
ACK received: 160  
Received ACK for START (write).  
Sent MID packet (write).  
ACK received: 162  
Received ACK for MID (write).  
Sent END packet (write) with checksum 0x0011892c  
ACK received: 164  
Got PACKET_TYPE_FUC  
FSM (write): Transfer encountered an error.  
  
==== Timing Summary ====  
Total round-trip time: 1680.06 ms  
Average response time: 559.52 ms over 3 measurements  
Per-packet response times (ms):  
    ACK_A0: 25.32 ms  
    ACK_A2: 177.97 ms  
    ACK_A3: 1475.28 ms
```

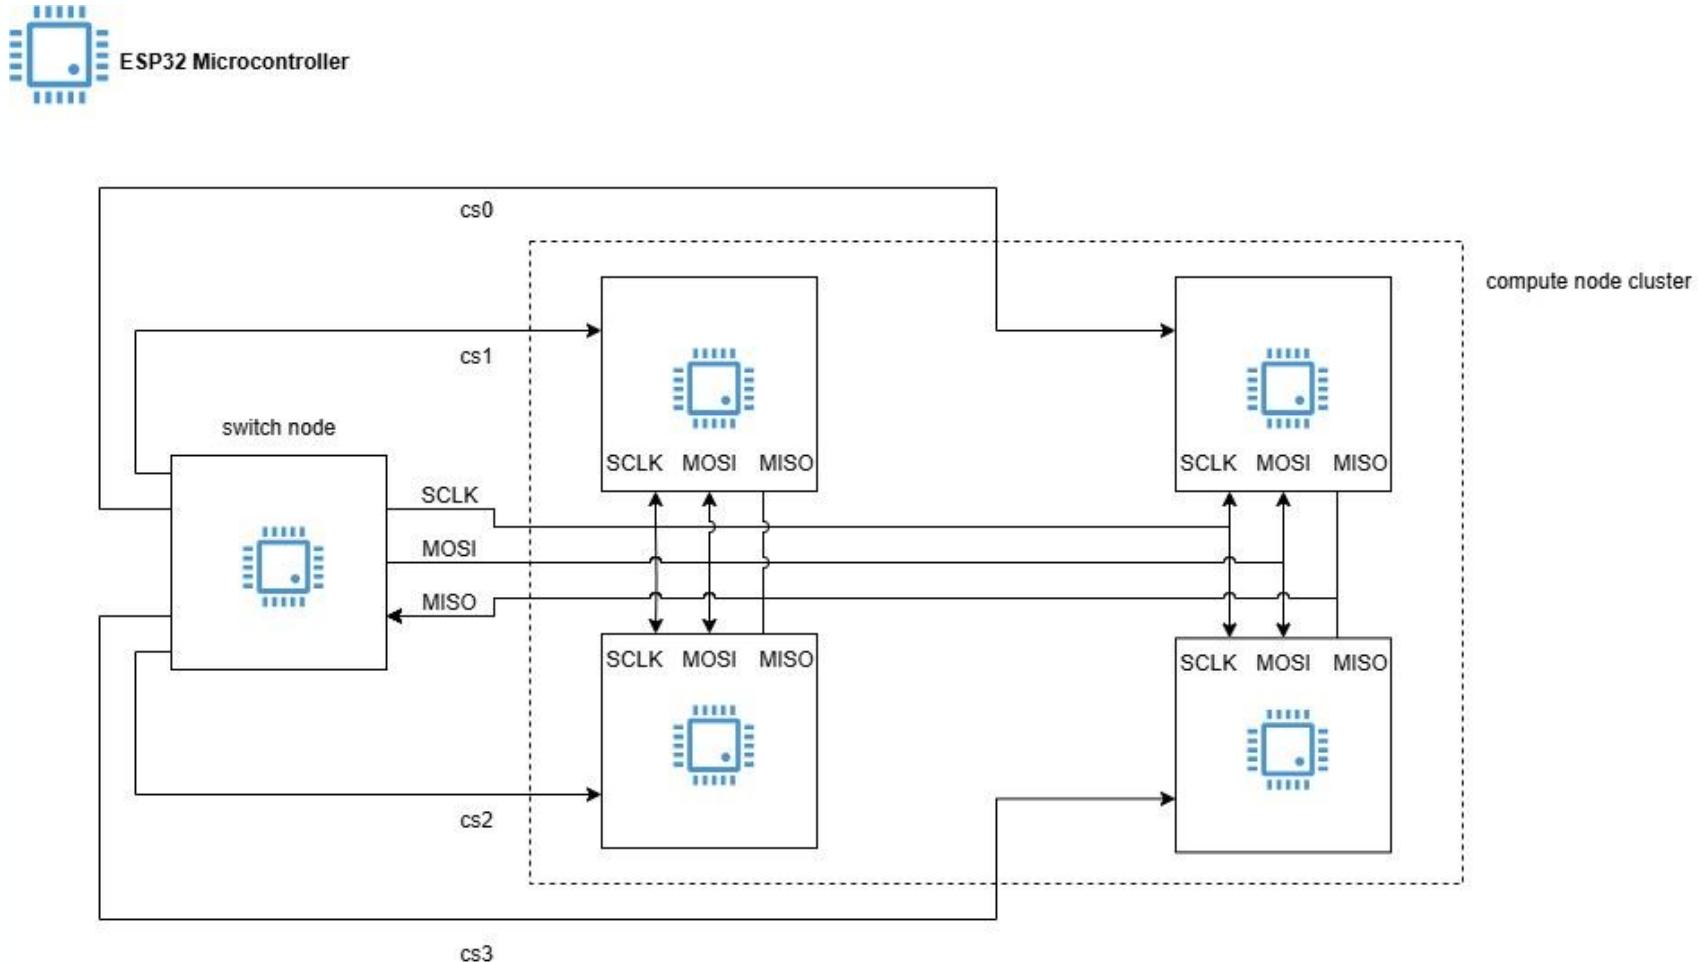
The Switch Nodes must be able to communicate errors to both the Network Node and Compute Nodes.



The Switch Nodes must maintain data when there is a power loss.

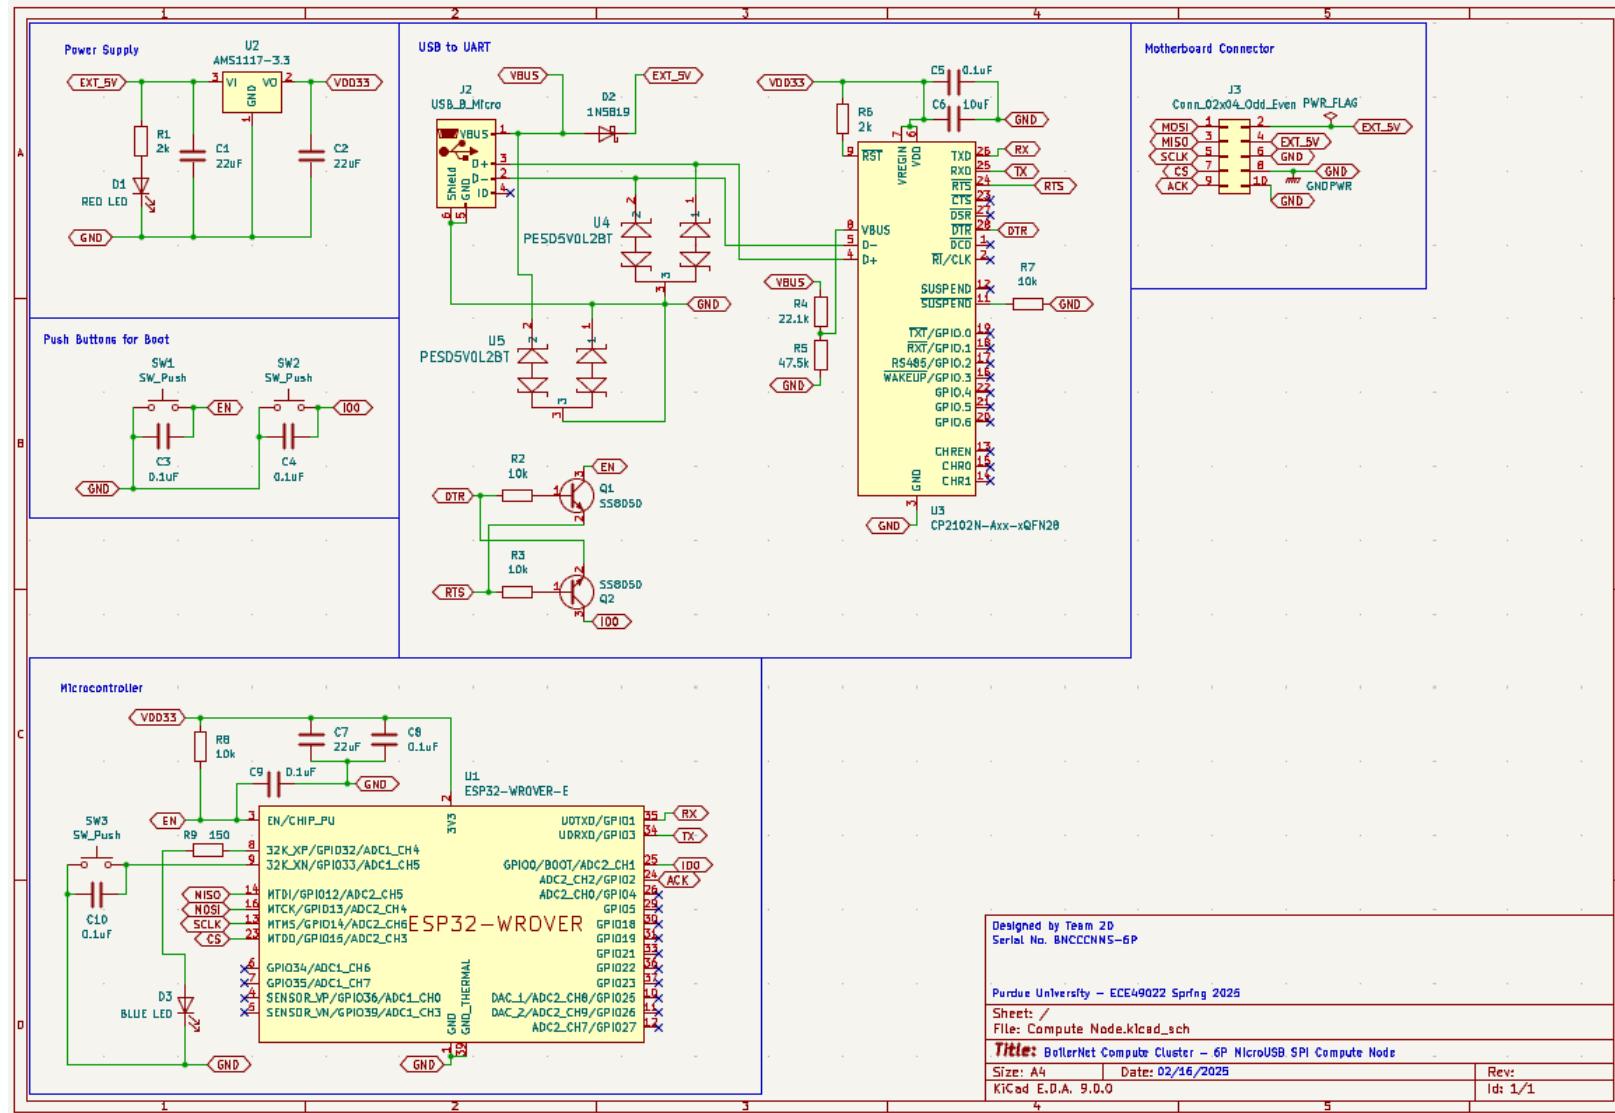
# Distributed Compute Layer

## Block Diagram



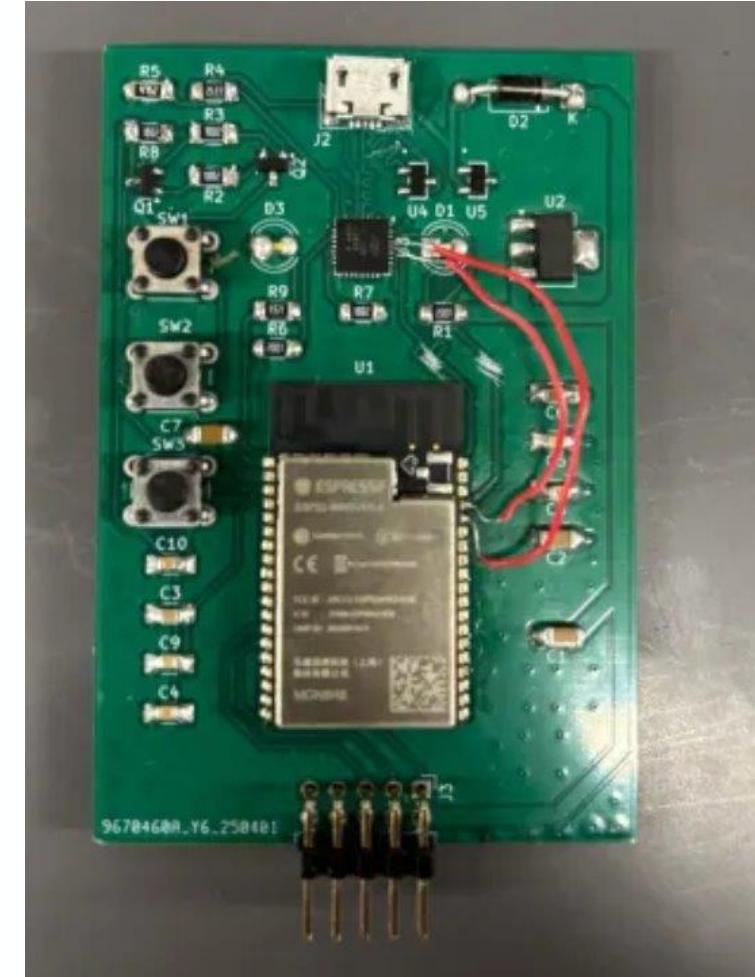
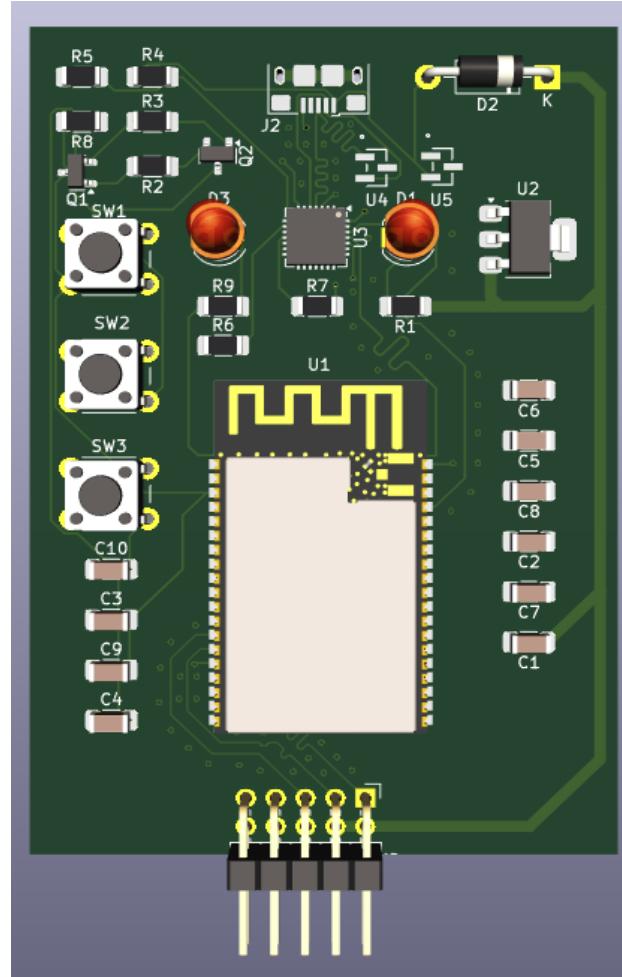
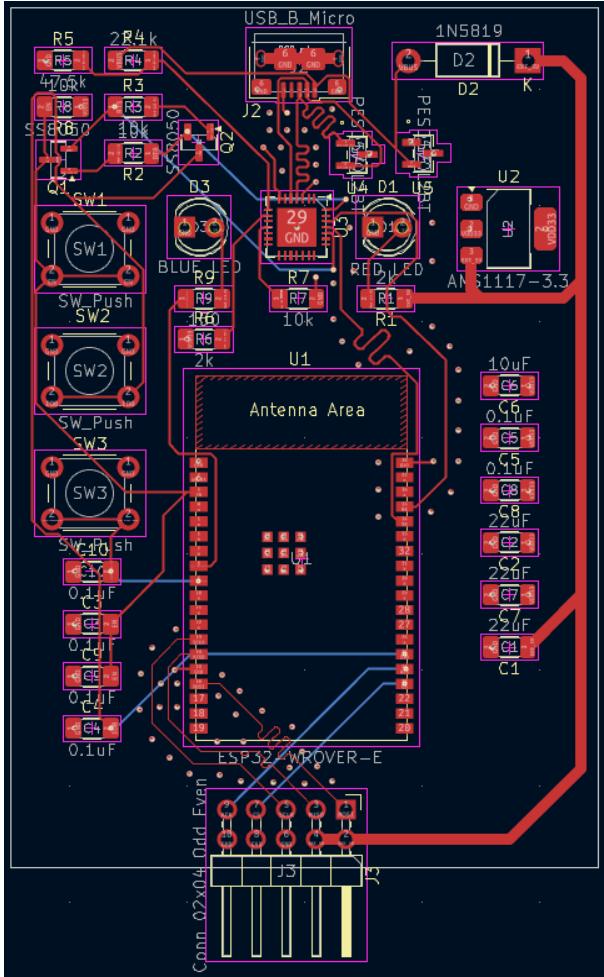
# Distributed Compute Layer

## Schematic



# Distributed Compute Layer

PCB

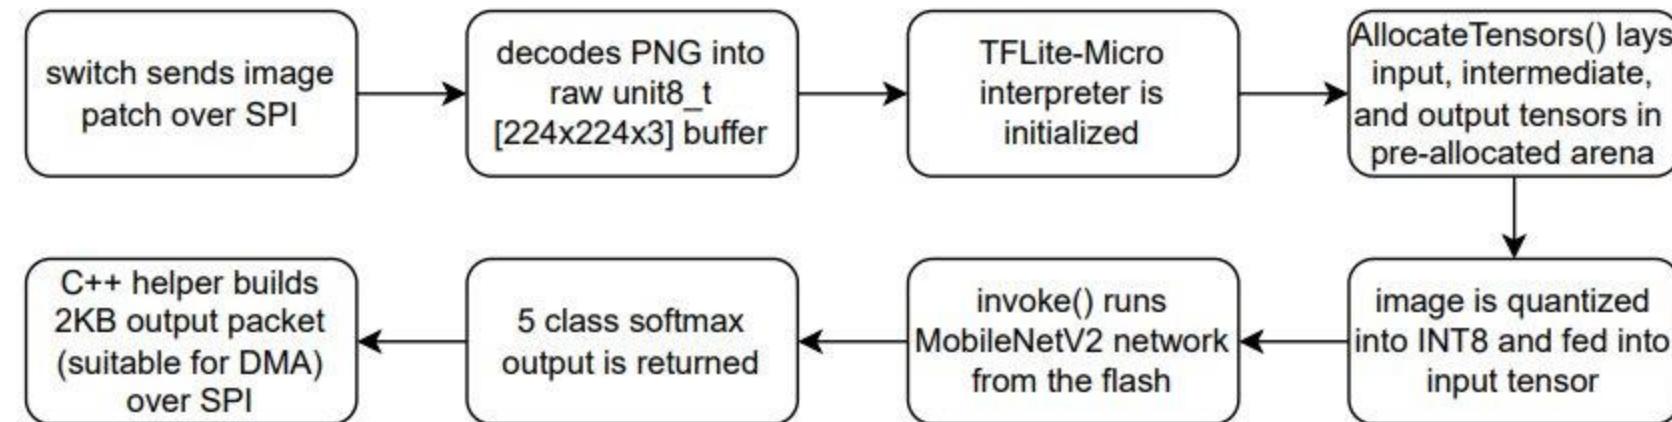


# Distributed Compute Layer

## Specifications

- Deep Learning Model running on a Compute Node must not take more than 10 seconds for inference
- Compute Nodes must decode PNG buffers into Byte Arrays in under 2s
- Compute Nodes must be able to fit the entire PNG Buffer, Quantized Activations and intermediate Buffers in RAM

## Activity Diagram



# *Distributed Compute Layer*

## Specification Relevance and Justification

- **Relevance:** Offloads workload away from switch nodes and parallelizes compute among 4 different compute nodes to speed up the overall image inference tasks
- **Justification:** Switch nodes already have multiple tasks such as interfacing with SD card and network nodes and adding computation to that would slow down the task completion speed.

## Specification Testability

- Speed of task completion can be tested by measuring the time it takes for the Deep Learning Model running on the Compute node for inference
- Size of task can be tested by checking size of PNG buffer sent without RAM overflow

## Specification Motivation

- The Distributed Compute Layer was devised in order to increase the speed of completing tasks to a reasonable level for the purposes of a compute cluster. The distributed compute layer aims to alleviate the issue of microprocessors having slower speeds by parallelizing the task

# Distributed Compute Layer

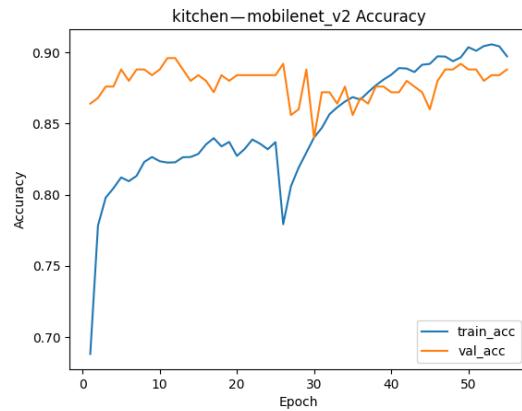
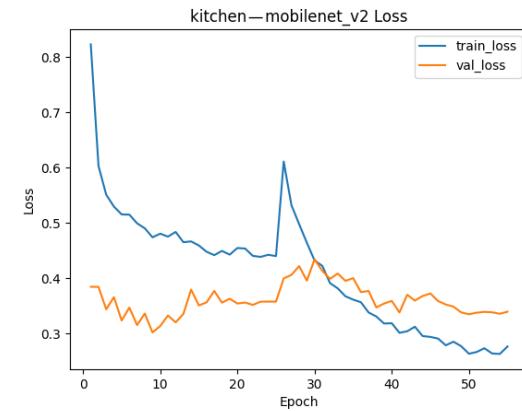
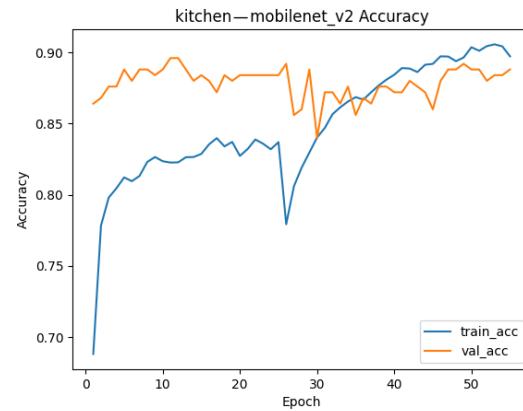
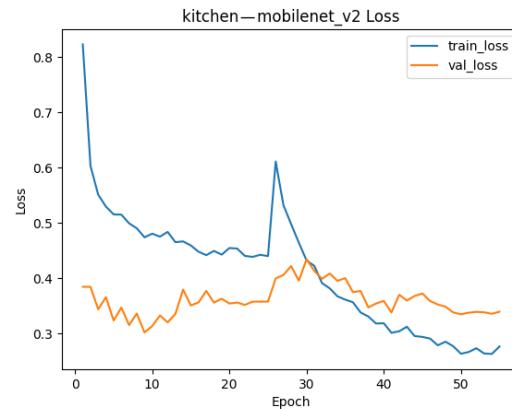
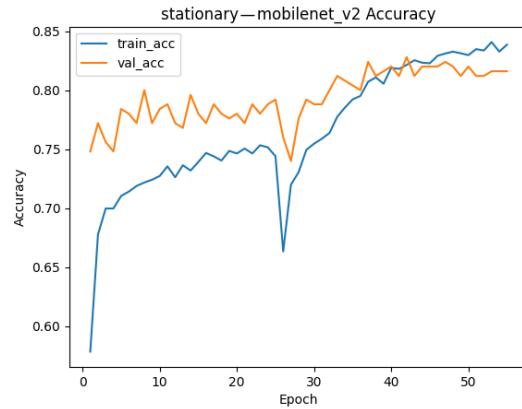
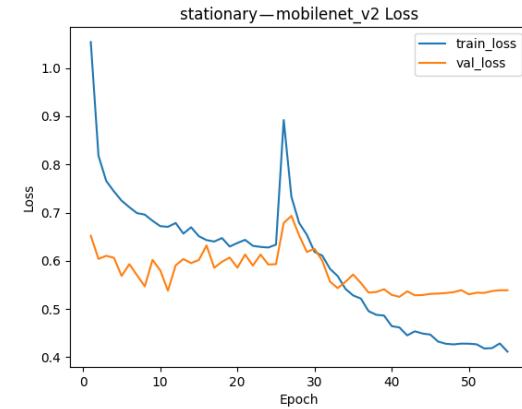
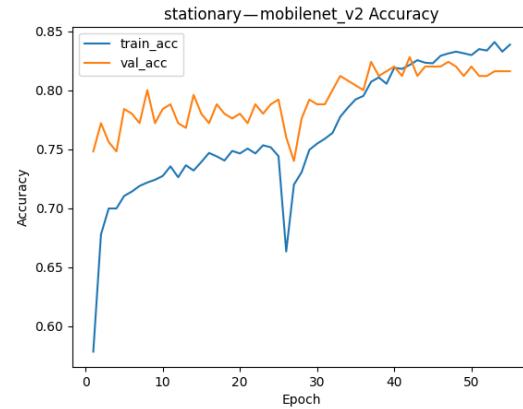
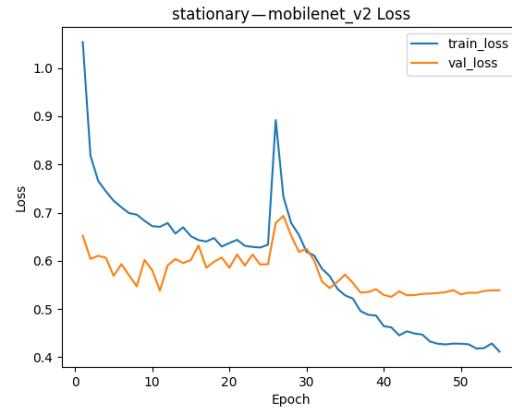
## Specification Evidence

- Deep Learning Model running on a Compute Node must not take more than 10 seconds for inference
  - Inference takes 9.017 seconds as shown by the 2nd yellow arrow
- Compute Nodes must decode PNG buffers into Byte Arrays in under 2s
  - LodePNG's open-source decoder and the WROVER's SPI-RAM allowed us to override general malloc() calls to help in larger working buffers resulting in a sub-500ms decoding time shown by the 2nd purple arrow
- Compute Nodes must be able to fit the entire PNG Buffer, Quantized Activations and intermediate Buffers in RAM
  - We utilize the SPIRAM-enabled WROVER IE boards to ensure internal RAM doesn't overflow. With the 224x224 tiles taking over 300KB and the TF-Lite's Tensor Arena requiring over 500KB of working memory, internal RAM would not be enough. We ensure robustness by cancelling transactions overflowing a preset utilization limit of 4MB. This is shown by the first purple arrow showing 4MB of PSRAM being initialized.

```
I (1471) esp_param: Adding pool of 4096K of PSRAM memory to heap allocator ←
I (1479) spi_flash: detected chip: generic
I (1481) spi_flash: flash io: dio
W (1484) spi_flash: Detected size(8192K) larger than the size in the binary image header(2048K). Using the size in the binary image header.
I (1498) main_task: Started on CPU0
I (1508) main_task: Calling app_main()
I (1508) gpio: GPIO[32]: InputEn: 0| OutputEn: 1| OpenDrain: 0| Pullup: 0| Pulldown: 0| Intr:0
I (1508) gpio: GPIO[2]: InputEn: 0| OutputEn: 1| OpenDrain: 0| Pullup: 0| Pulldown: 0| Intr:0
I (1518) BLADE_SPI: BLADE HANDLER INITIALIZED!
I (1518) main_task: Returned from app_main()
I (1538) BLADE_SPI: Max internal used: 3342336
I (1658) BLADE_SPI: 0: Sending buffer; first bytes: 0x00 0x00 0x00 0x00 ... 0x00 0x00
I (20298) BLADE_SPI: 0: Received buffer; first bytes: 0x00 0x00 0x00 0x00 ... 0x00 0x00
I (20298) BLADE_SPI: IN BLADE_FSM_WAIT_START
I (20298) BLADE_SPI: START_DATA_TO_BLADE
I (20308) BLADE_SPI: BLADE_FSM_WAIT_START, starting receive: internal free=3667676, psram free=3372812
I (20318) BLADE_SPI: align=0
W (20318) BLADE_SPI: PNGBUF made!
I (22808) BLADE_SPI: IN BLADE_FSM RECEIVE
I (22818) BLADE_SPI: In RECEIVE: Got 2038 bytes
I (22828) BLADE_SPI: 1: Sending buffer; first bytes: 0xD1 0x07 0xF6 0x0B ... 0x00 0x00
I (22988) BLADE_SPI: 1: Received buffer; first bytes: 0xD1 0x07 0xF6 0xA3 ... 0x00 0x00
I (22988) BLADE_SPI: IN BLADE_FSM RECEIVE
I (22988) BLADE_SPI: In RECEIVE: Got 2038 bytes
I (22998) BLADE_SPI: 0: Sending buffer; first bytes: 0xD1 0x07 0xF6 0xA3 ... 0x00 0x00
I (22998) BLADE_SPI: 0: Received buffer; first bytes: 0xD2 0x07 0xB2 0x87 ... 0x00 0x00
I (23008) BLADE_SPI: IN BLADE_FSM RECEIVE
I (23008) BLADE_SPI: In RECEIVE: Got 1970 bytes
I (23018) BLADE_SPI: SPI recv 89604 bytes in 2695 ms
I (23018) BLADE_SPI: In BLADE_FSM_PROCESS; BEGINNING PROCESSING! →
I (23038) BLADE_SPI: HEADER: 137, 80, 78, 71, 13, 10, 26, 10
I (23448) BLADE_SPI: PNG decode in 415 ms ←
I (32508) BLADE_SPI: Inference in 9017 ms ←
I (32508) BLADE_SPI: Result: Egyptian ←
I (32528) BLADE_SPI: 1: Sending buffer; first bytes: 0xD3 0x00 0x80 0x7B ... 0x00 0x00
I (32528) BLADE_SPI: 1: Received buffer; first bytes: 0xD3 0x07 0xB2 0x87 ... 0x00 0x00
I (32528) BLADE_SPI: IN BLADE_FSM_WAIT_START
I (32538) BLADE_SPI: 0: Sending buffer; first bytes: 0x00 0x00 0x00 0x00 ... 0x00 0x00
```

# Distributed Compute Layer

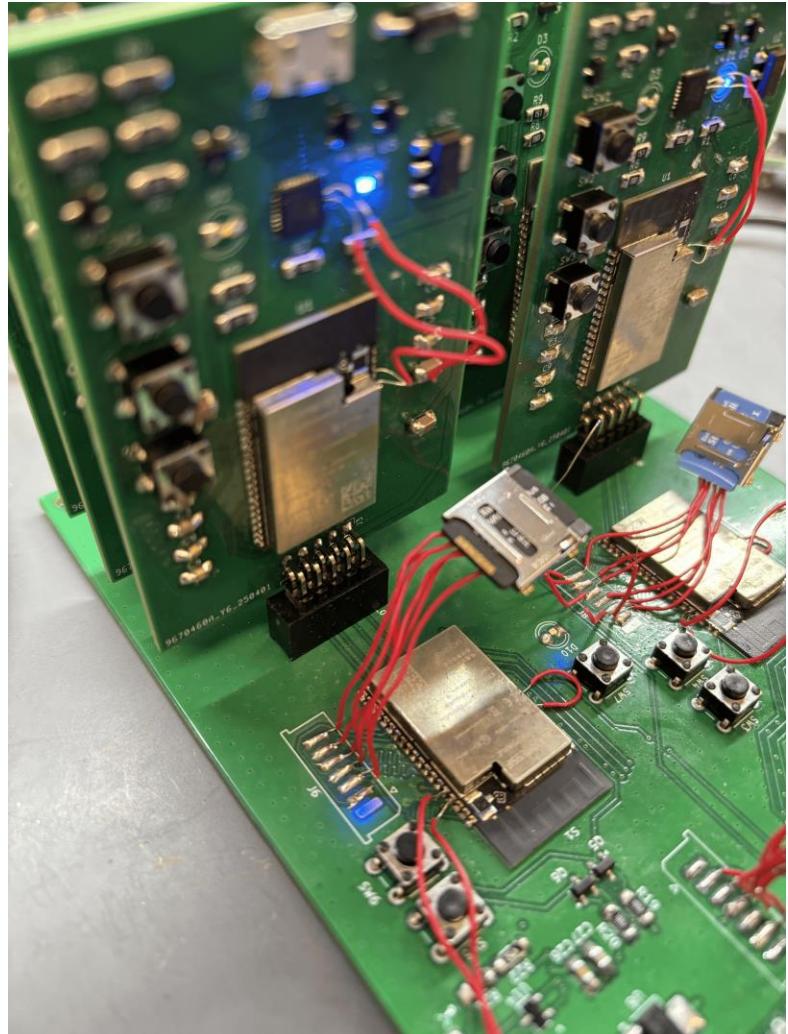
## Specification Evidence



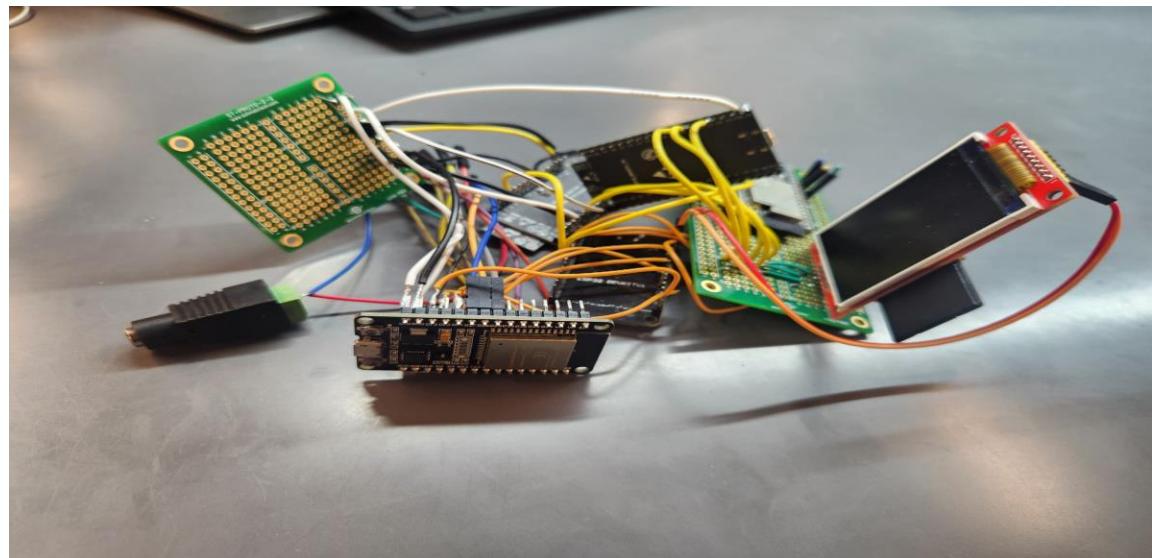
# In-Retrospect!

- Footprint PCB Management
  - Barrel Jack direction
  - Sd Card Reader footprint
  - Capacitor and Resistor Packages
- RX TX trace allignment
- GPIO Matrix Datasheet Awareness

Name	No.	Type	Function
SENSOR_VP	4	I	GPIO36, ADC1_CH0, RTC_GPIO0
SENSOR_VN	5	I	GPIO39, ADC1_CH3, RTC_GPIO3
IO34	6	I	GPIO34, ADC1_CH6, RTC_GPIO4
IO35	7	I	GPIO35, ADC1_CH7, RTC_GPIO5
IO32	8	I/O	GPIO32, XTAL_32K_P (32.768 kHz crystal oscillator input), ADC1_CH4, TOUCH9, RTC_GPIO9
IO33	9	I/O	GPIO33, XTAL_32K_N (32.768 kHz crystal oscillator output), ADC1_CH5, TOUCH8, RTC_GPIO8
IO25	10	I/O	GPIO25, DAC_1, ADC2_CH8, RTC_GPIO6, EMAC_RXD0
IO26	11	I/O	GPIO26, DAC_2, ADC2_CH9, RTC_GPIO7, EMAC_RXD1



# *Project Progression!*



# *Thank You*

Q&A

